

Progress of the Project

Tsung-Min Pai

2023/9/7

Outline

- **GAT**
 - Experiment 1 and 2
 - Experiment 3
- **Future Work**

Graph Attention Network - GAT

Graph Attention Network - GAT

Model:

```
class GAT(nn.Module):
    def __init__(self, in_dim, hidden_dim, out_dim, num_heads, dropout_prob=0.2):
        super(GAT, self).__init__()
        # do not check the zero in_degree since we have all the complete graph
        self.layer1 = GATConv(in_dim, hidden_dim, num_heads=num_heads, activation=F.relu, allow_zero_in_degree=True)
        self.layer2 = GATConv(hidden_dim * num_heads, out_dim, num_heads=num_heads, allow_zero_in_degree=True)

        # Adding Dropout for regularization
        self.dropout = nn.Dropout(dropout_prob)

    def forward(self, g, h):
        # Apply GAT layers
        h = self.layer1(g, h)
        h = h.view(h.shape[0], -1)
        h = F.relu(h)
        h = self.dropout(h)
        h = self.layer2(g, h).squeeze(1)

        # Store the output as a new node feature
        g.ndata['h_out'] = h


        # Use mean pooling to aggregate this new node feature
        h_agg = dgl.mean_nodes(g, feat='h_out')
        return h_agg
```

- Use the **new** version of the dataset

Experiment 1 and 2

Experiment 1 and 2

- **Experiment 1:**
 - Dataset is 165 APs with 11 versions of embedding
 - Graph classification
- **Experiment 2:**
 - Experiment 1 + **benign** data
 - Benign made from benign.txt → 1000 graphs
 - Graph classification



Dataset

Format:

```
{"label": 10, "num_nodes": 3, "node_feat": [205565, 733769, 250773], "edge_attr": [23, 23], "edge_index": [[0, 0], [1, 2]]}  
{"label": 11, "num_nodes": 3, "node_feat": [470650, 663446, 627322], "edge_attr": [23, 23], "edge_index": [[0, 0], [1, 2]]}  
{"label": 15, "num_nodes": 2, "node_feat": [9863, 103498], "edge_attr": [23], "edge_index": [[0], [1]]}  
{"label": 16, "num_nodes": 2, "node_feat": [157277, 753159], "edge_attr": [23], "edge_index": [[0], [1]]}  
{"label": 22, "num_nodes": 36, "node_feat": [83068, 614681, 444724, 266227, 121794, 623948, 116790, 769462, 255741, 169794,
```


- Have 165 APs, each AP has 1000 variation → nodes are different but relations are same
- 0~99 test, 100~199 validation, 200~999 train → 1:1:8
- Use transR_50, transE_50, transH_50, secureBERT... as embedding → **11 versions**
- **secureBERT** → dimension = 250, 150, 100, 50

```
pca = PCA(n_components=DIM)  
ent_embeddings = pca.fit_transform(ent_embeddings)
```


Dataset

- Benign graphs for experiment 2:

1. 400: Leaf nodes + their source nodes




2. 300: Leaf nodes + their source nodes with source nodes' neighbor nodes



3. 200: Leaf nodes + their 2 layer source nodes



4. 100: Leaf nodes + their 2 layer source nodes with source nodes' neighbor nodes



- For version 2 and 4 graphs:

- Constrain the # of relation between the same nodes within 8
- Constrain the # of the triplets in the graph within 32

Experiment 1 and 2

- Record the training in a **log file**
- Also record the **classification report** supported by sklearn
- Give these files to Euni
- Log file:

```
08/29/2023, 09:04:02# labels of Validation: tensor([ 96,  4, 129, 111,  84, 162, 102, 103, 132, 114,  32, 110,  57,  85,
          149,  74], device='cuda:3') torch.Size([16])
08/29/2023, 09:04:02# predicted of Validation: tensor([ 96, 153, 113, 153,  84, 153, 153, 113, 161, 113,  32, 113, 153,  85,
          149, 153], device='cuda:3') torch.Size([16])
08/29/2023, 09:04:03# Validation Loss: 2.6715 | Validation Accuracy: 0.3879
```

- Classification report:

	precision	recall	f1-score	support
T1003.001_0ef4cc7b-611c-4237-b20b-db36b6906554	1.00	1.00	1.00	100
T1003.001_35d92515122effdd73801c6ac3021da7	1.00	1.00	1.00	100
T1003.002_5a484b65c247675e3b7ada4ba648d376	0.00	0.00	0.00	100
T1003.002_7fa4ea18694f2552547b65e23952cabb	1.00	1.00	1.00	100
T1003.003_9f73269695e54311dd61dc68940fb3e1	0.00	0.00	0.00	100
T1003.003_f049b89533298c2d6cd37a940248b219	0.00	0.00	0.00	100

Similar with the MLP, RNN:
More triplets, more accurate

Experiment 1 and 2

- Total: 25 epochs
 - Optimizer = AdamW(model.parameters(), lr=5e-4)
 - Criterion = nn.CrossEntropyLoss()
 - Batch size = 4
- Except **secureBERT**, all about 35~40% test accuracy
- secureBERT family \approx 12% test accuracy
- Experiment 1 and 2 have similar performance
 - since benign only has 1000 graph \rightarrow kind of balance
 - If we make the benign graph much more than AP(real data) \rightarrow imbalance

Experiment 3


Experiment 3

- **Experiment 3:**
 - Consider the **neighbor** benign nodes
 - Edge classification
 - I think it's more like an triplet classification?
- Given a graph → label the triplets with the benign or the specific AP

Source txt file:

```
853776 595218 13 a
593289 563219 17 b
388326 563219 17 b
```


- a means attack pattern
- b means benign



Dataset

Format in experiment 1 and 2:

```
{"label": 10, "num_nodes": 3, "node_feat": [205565, 733769, 250773], "edge_attr": [23, 23], "edge_index": [[0, 0], [1, 2]]}  
{"label": 11, "num_nodes": 3, "node_feat": [470650, 663446, 627322], "edge_attr": [23, 23], "edge_index": [[0, 0], [1, 2]]}  
{"label": 15, "num_nodes": 2, "node_feat": [9863, 103498], "edge_attr": [23], "edge_index": [[0], [1]]}  
{"label": 16, "num_nodes": 2, "node_feat": [157277, 753159], "edge_attr": [23], "edge_index": [[0], [1]]}  
{"label": 22, "num_nodes": 36, "node_feat": [83068, 614681, 444724, 266227, 121794, 623948, 116790, 769462, 255741, 169794,
```



Format in experiment 3:

```
{"labels": [45, 65, 45, 45], "num_nodes": 4, "node_feat": [578353, 695633, 234474, 883199], "edge_attr": [24, 2, 7, 2],  
{"labels": [45, 65, 45, 45], "num_nodes": 4, "node_feat": [578353, 234474, 1085219, 1079260], "edge_attr": [24, 2, 7, 2]  
{"labels": [45, 65, 45, 45], "num_nodes": 4, "node_feat": [578353, 946954, 234474, 391415], "edge_attr": [24, 2, 7, 2],
```

- From graph classification to **edge classification**, which is **multi-label** classification
- # of labels = # of triplets
- Haven't successfully trained → some tensor error

Future Work

Future Work

- **GAT**

- Successfully run the experiment 3
- Try GCN or different architecture of the model
- Improve the performance of the model (if available)

Thanks!!

Appendix

Useful Links

<https://zhuanlan.zhihu.com/p/107737824>

[https://blog.csdn.net/uncle ll/article/details/82778750](https://blog.csdn.net/uncle_ll/article/details/82778750)

[https://docs.dgl.ai/en/1.1.x/guide cn/minibatch-edge.html#guide-cn-minibatch-edge-classification-sampler](https://docs.dgl.ai/en/1.1.x/guide_cn/minibatch-edge.html#guide-cn-minibatch-edge-classification-sampler)

Experiment 3

```
binary_labels = torch.zeros(batch_size, num_classes)
for i, sample_labels in enumerate(labels):
    for label in sample_labels:
        binary_labels[i, label] = 1
```

```
preds = (logits >= 0.5).to(torch.float32)
```

```
accuracy = torch.mean((preds == binary_labels).float())
```

```
criterion = nn.BCEWithLogitsLoss()
```

- Total: 25 epochs
 - Optimizer = AdamW(model.parameters(), lr=5e-4)
 - Criterion = nn.CrossEntropyLoss()
 - Batch size = 4

Graph Convolutional Network - GCN

Graph Convolutional Network - GCN

Model:

```
class GCN(nn.Module):
    def __init__(self, in_feats, hidden_size, num_classes):
        super(GCN, self).__init__()
        self.conv1 = GraphConv(in_feats, hidden_size)
        self.conv2 = GraphConv(hidden_size, num_classes)

    def forward(self, g, inputs):
        h = self.conv1(g, inputs)
        h = torch.relu(h)
        h = self.conv2(g, h)

        g.ndata['h'] = h
        hg = dgl.mean_nodes(g, 'h')
        return hg
```

- Use the **old** version of the dataset
- Use **DGL** to be our library
- DGL data format:

```
batched_g is like:
Graph(num_nodes=96, num_edges=160, ndata_schemes={'feat': Scheme(shape=(1,), dtype=torch.int64)}, edata_schemes={})
num_nodes = 3*batch_size, num_edges = 5*batch_size
```

```
labels is like: tensor([ 76,  0,  0,  0,  0,  0,  0,  0,  0, 76,  0, 76,  0,  0,
                        0,  0, 76,  0, 30, 92,  0,  0, 76,  0,  0,  0,  0,
                        116,  0, 76, 76])
```

Result:

```
0%|          | 0/120 [00:00<?, ?it/s]
Epoch 0 | Train Loss: 2625.5943 | Train Accuracy: 0.4763
1%|          | 1/120 [00:56<1:52:21, 56.65s/it]
Validation Loss: 494.0275 | Validation Accuracy: 0.6642
99%|██████████| 119/120 [1:51:06<00:55, 55.13s/it]
Validation Loss: 0.9964 | Validation Accuracy: 0.6642
Epoch 119 | Train Loss: 0.9625 | Train Accuracy: 0.6644
100%|██████████| 120/120 [1:52:03<00:00, 56.03s/it]
Validation Loss: 0.9965 | Validation Accuracy: 0.6642
```

Test Accuracy: 66 %

- GAT applied on the old data has the similar result

Appendix

