

Progress of the Project

Tsung-Min Pai

2023/11/17

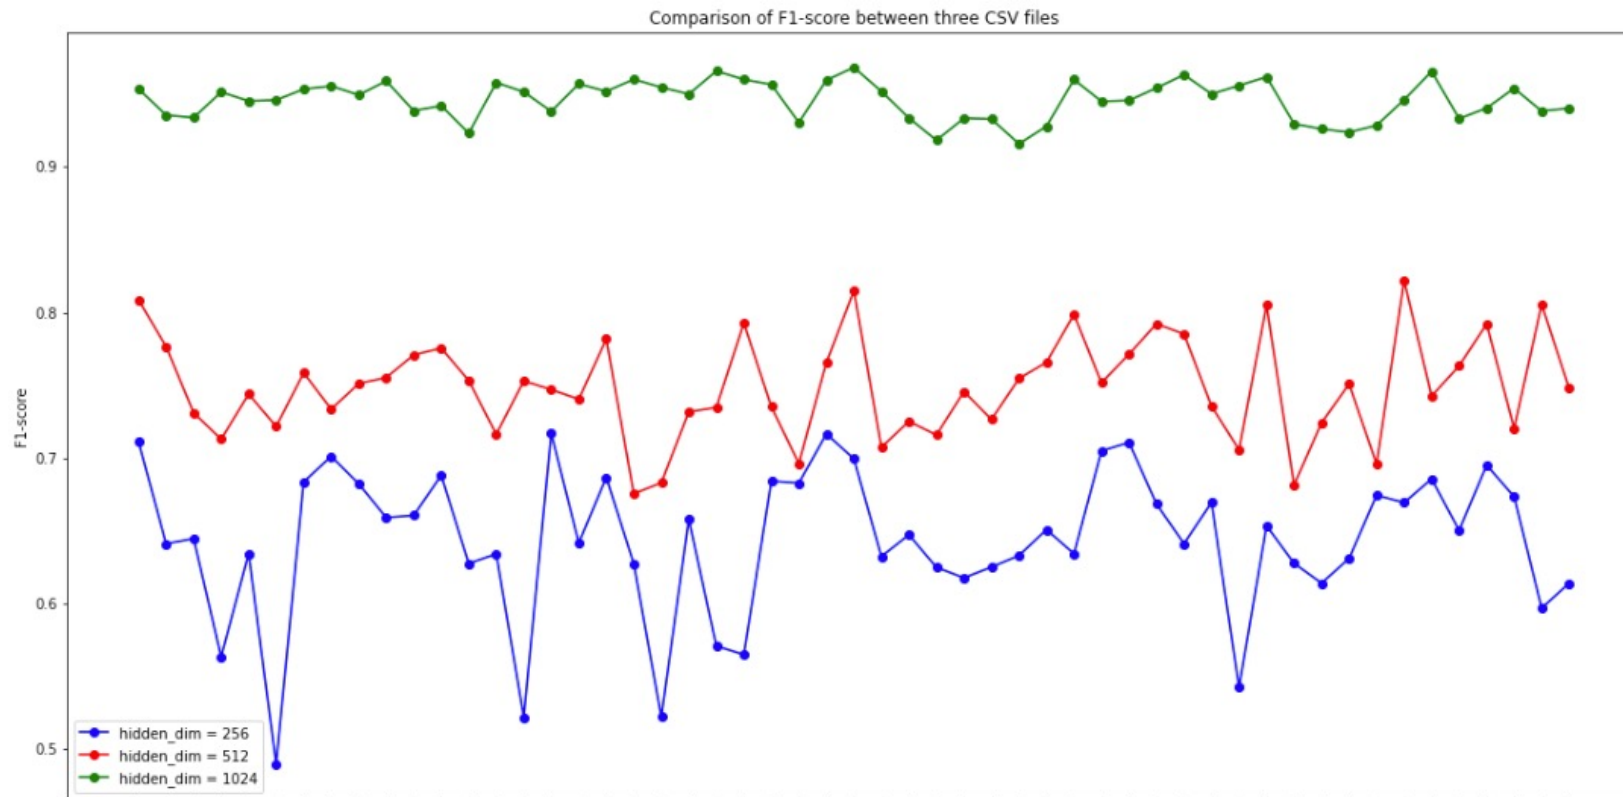
Outline

- **GNN**
 - **Experiment**
 - Recap
 - Observe the distribution of the prediction
 - Remove the predictable TTP from the dataset
 - **Paper**
 - GraphSMOTE
- **Future Work**

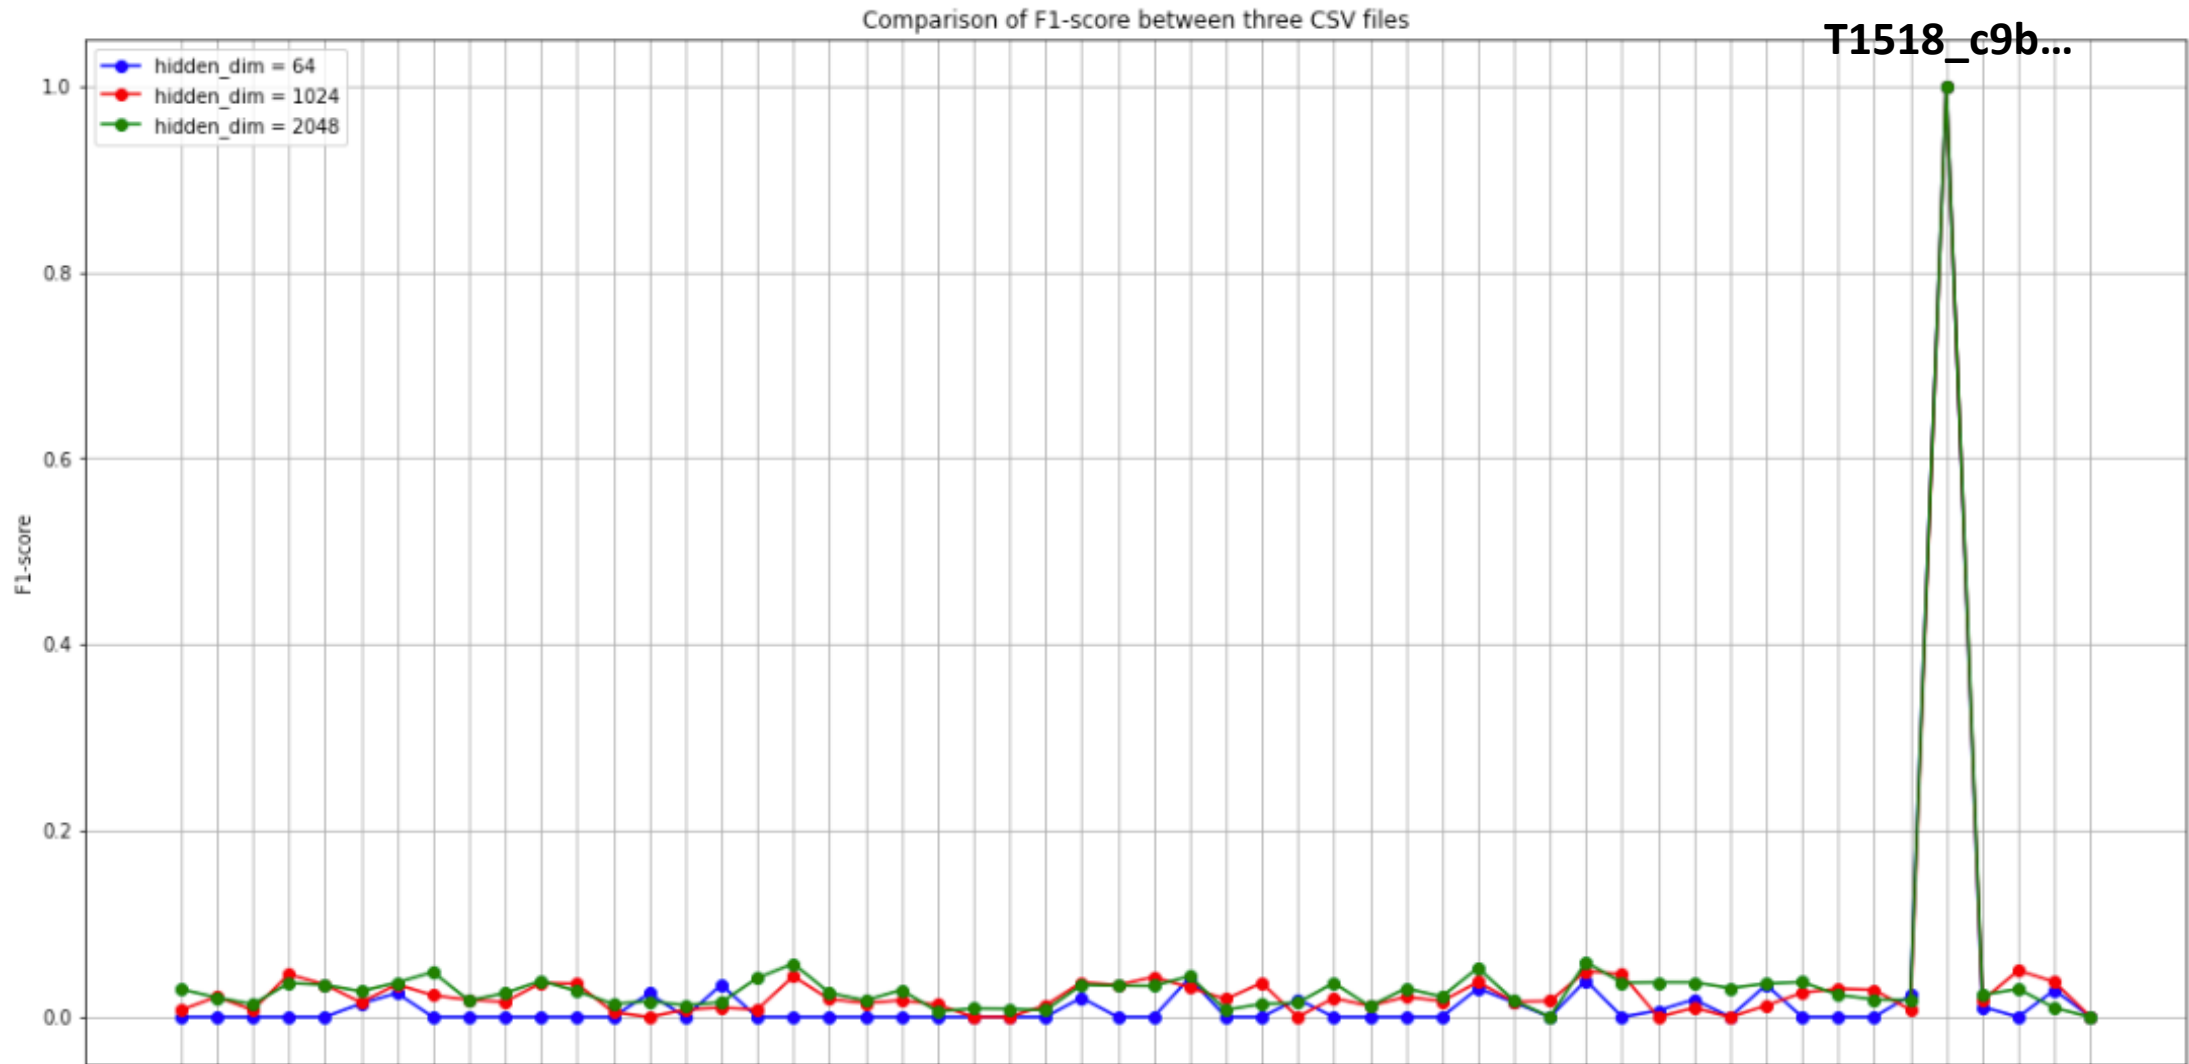
Experiment - Recap

Oversampling

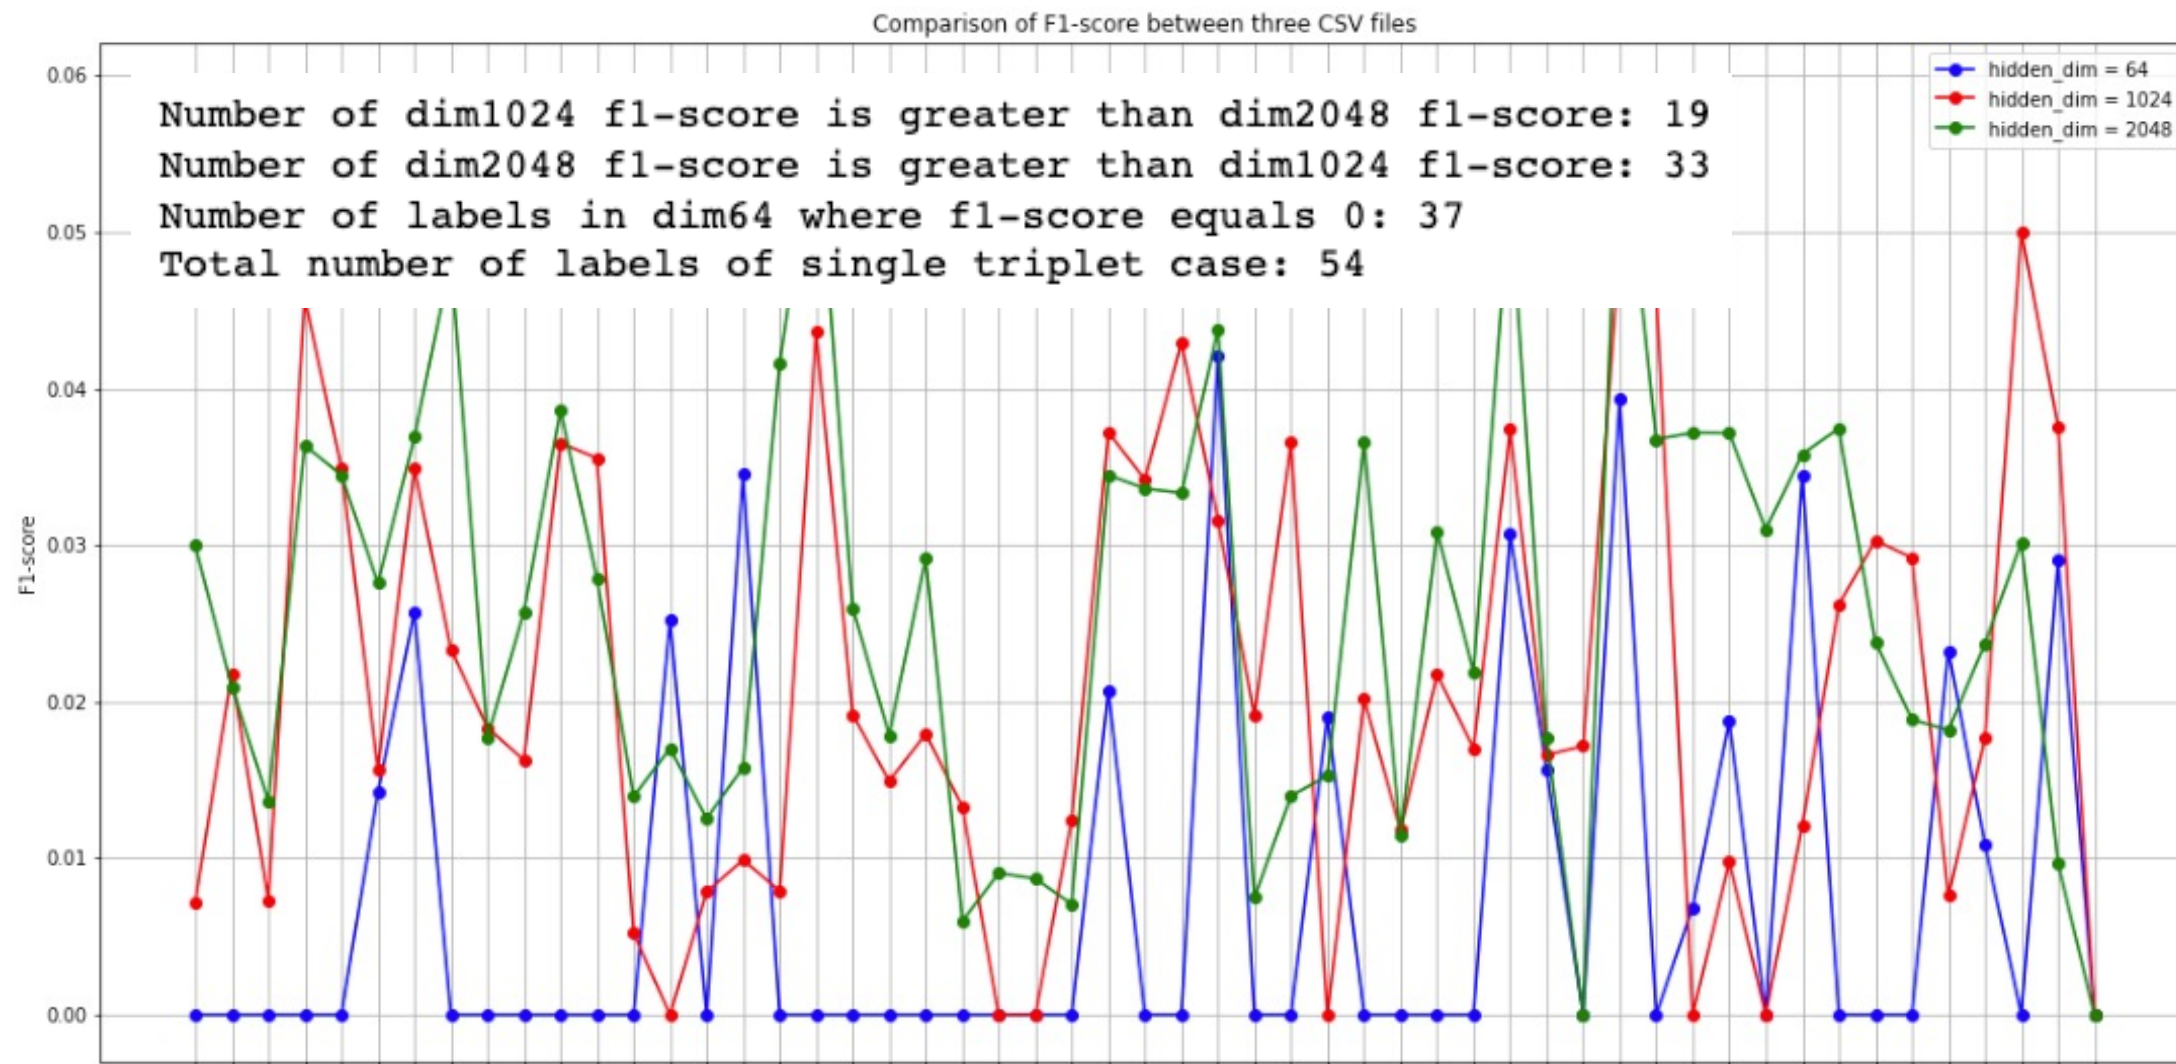
- Previous Trial: hope to see the result on **training** dataset
 - Use data with **320** times single triplet → # of training data = 13657600
 - Larger hidden **dimension** → more neurons to remember the data
 - Let the model **overfit** first → Succeed



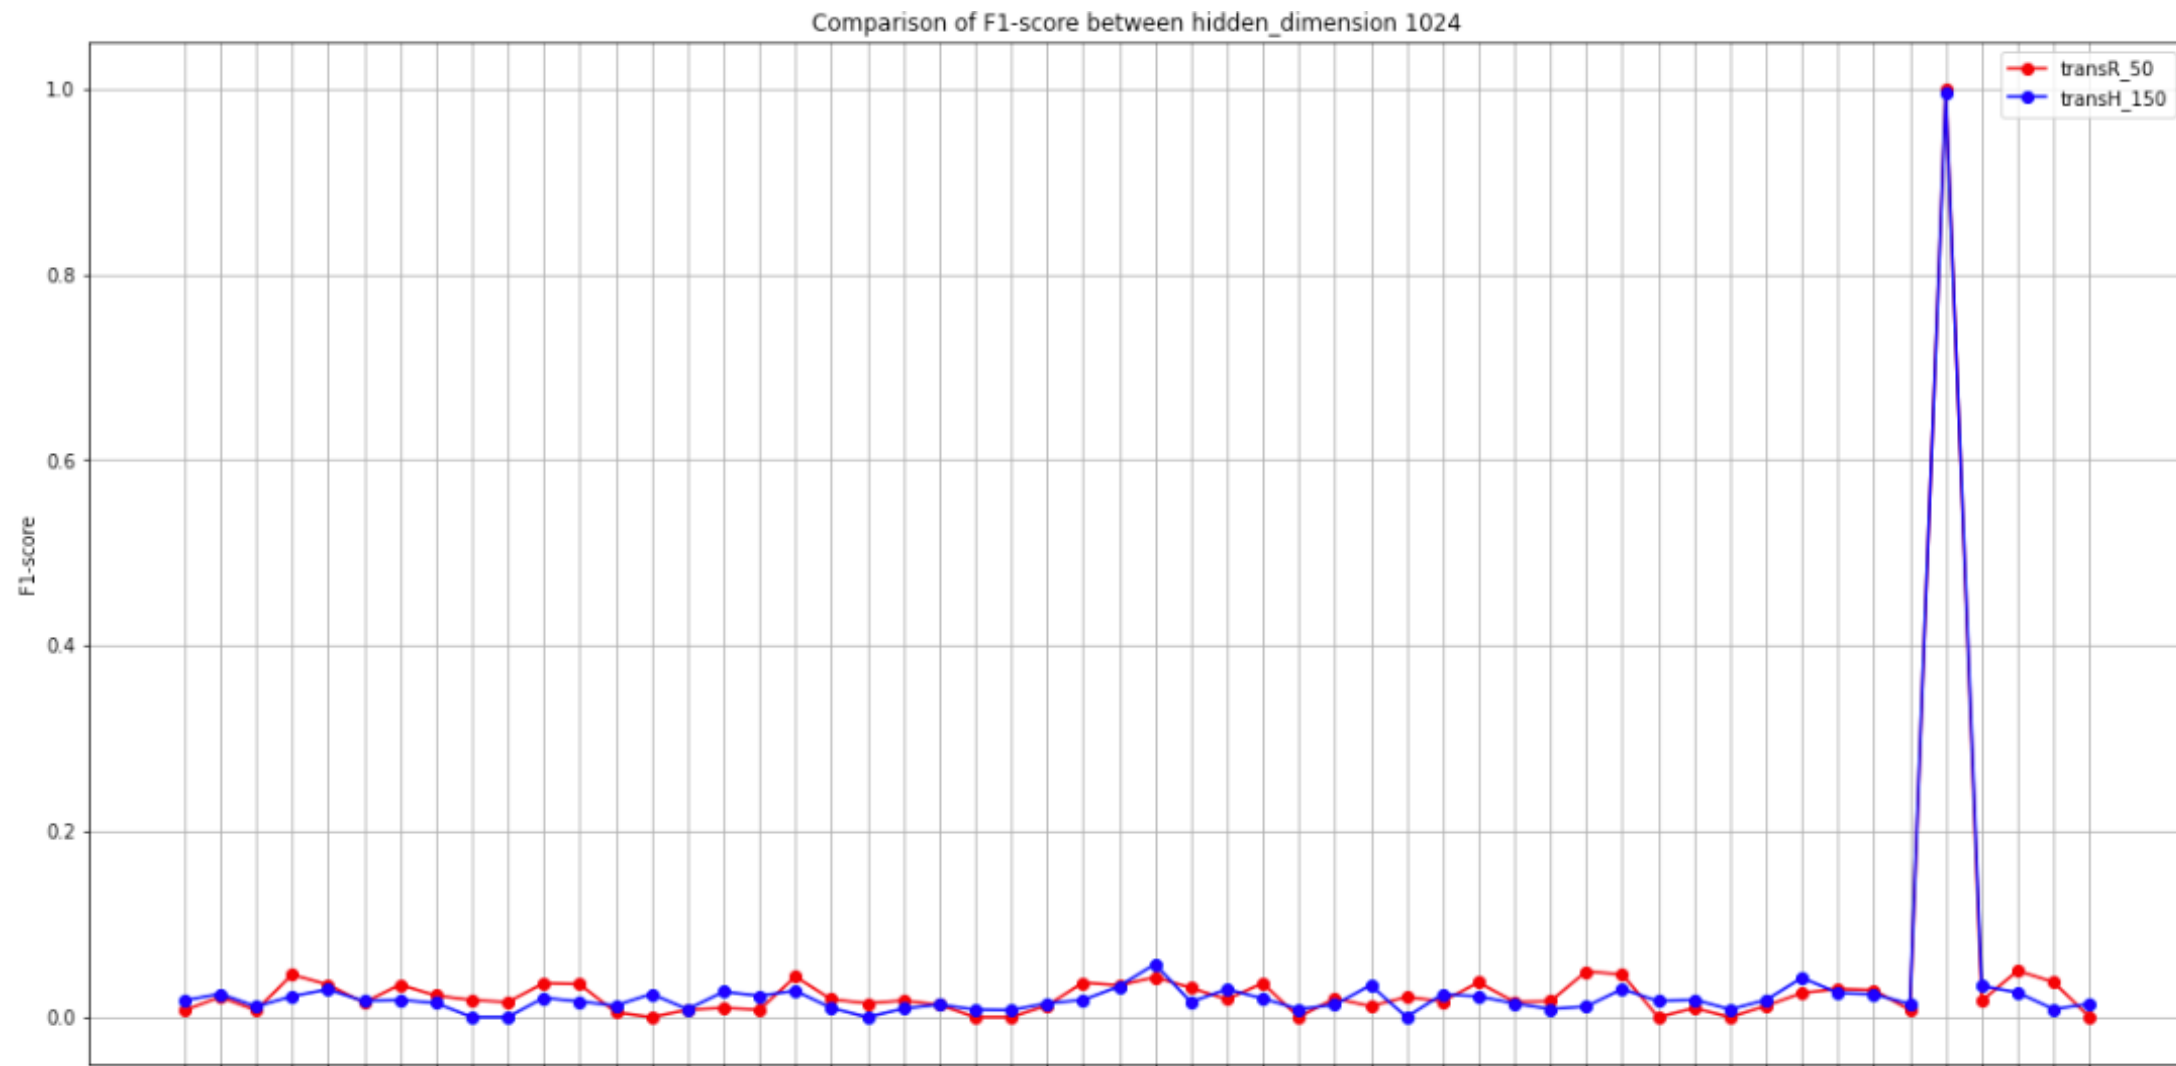
Observation on Different Dimension



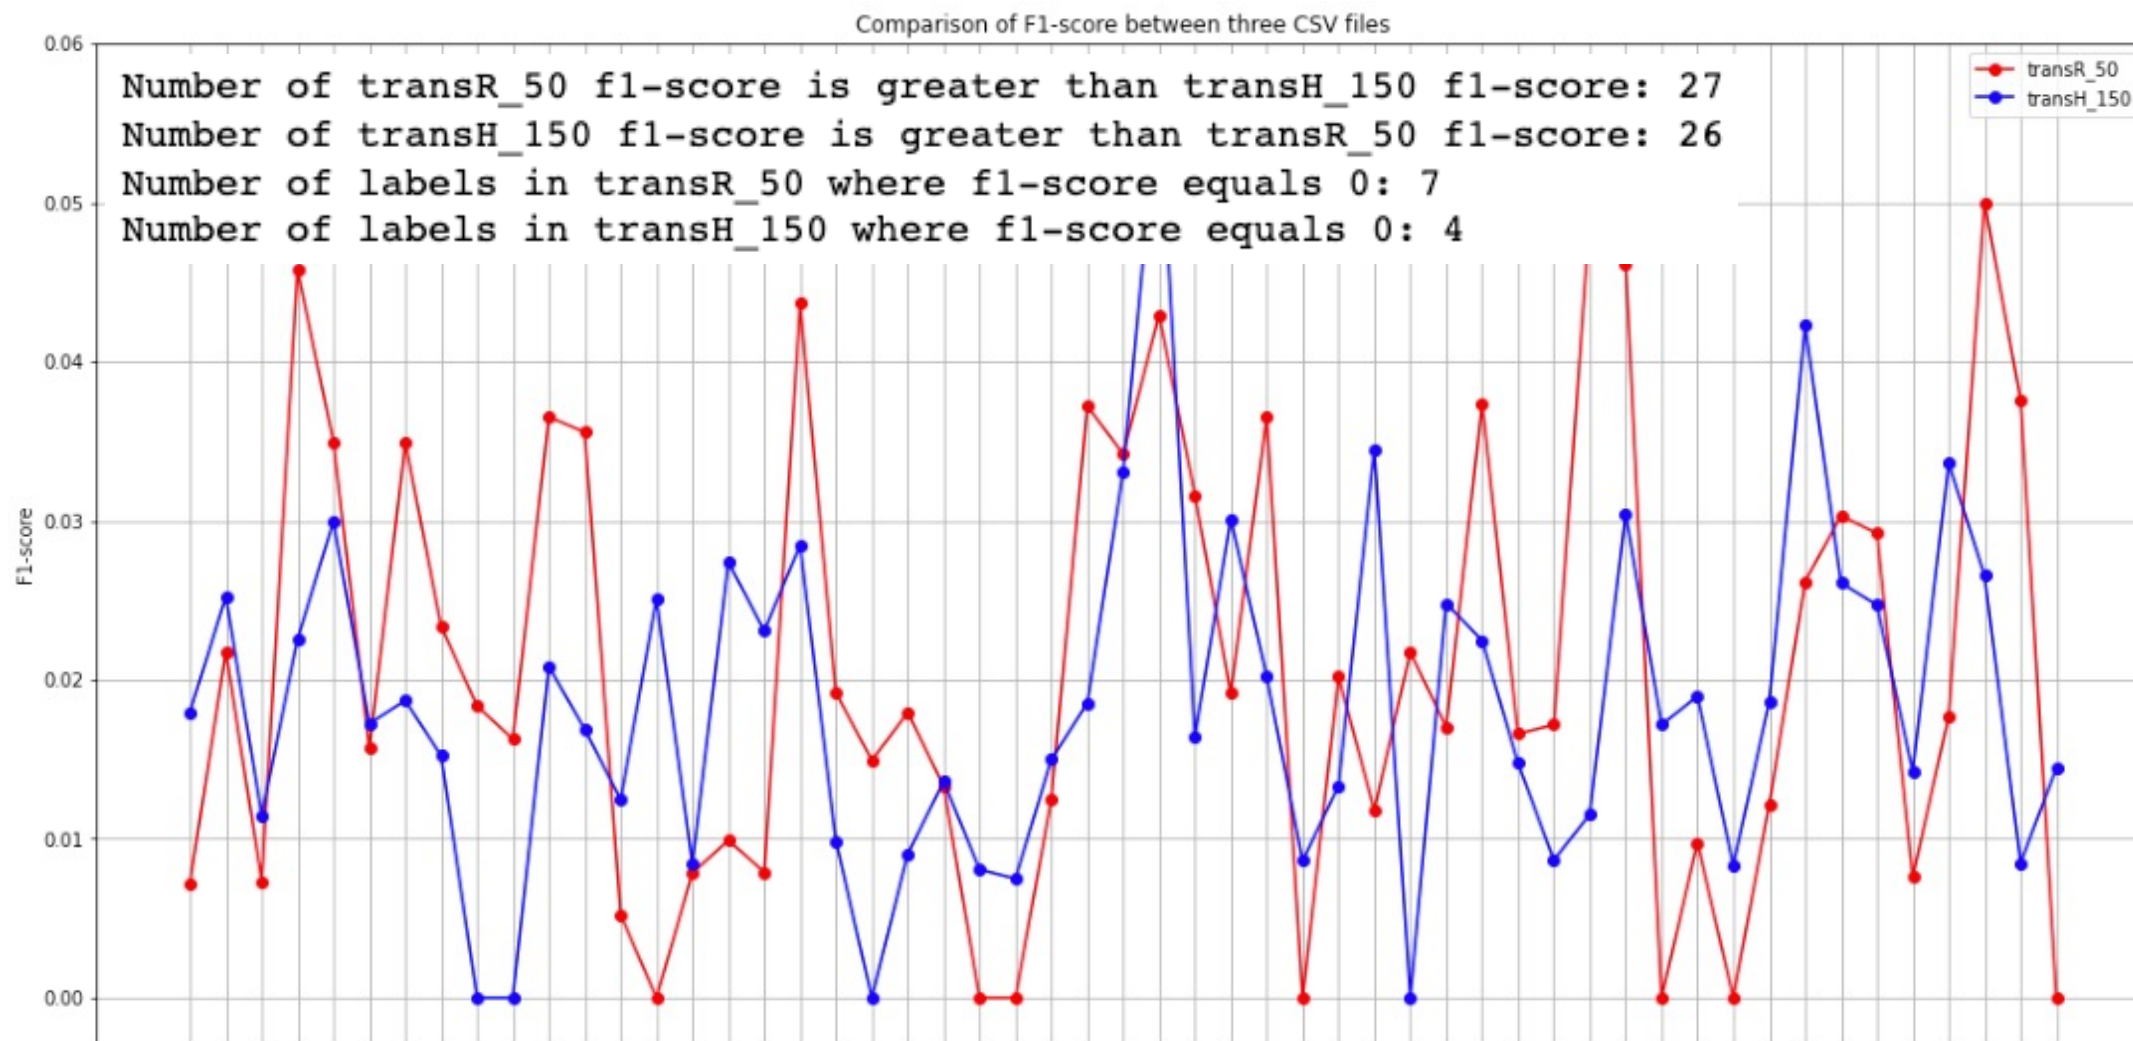
Observation on Different Dimension



Observation on Different Embedding



Observation on Different Embedding



Conclusion

- Noticed that **T1518_c9b** always got predicted in all the experiments
- **Hidden Dimension** do have an effect on the result → but it's all about from 0 to 0.05
- **Embedding** (transR_50 and tansH_150) seems to have the similar result → try more

Experiment - Observe

Thoughts

- Try the ensemble
 - Different GNN
 - Different **hidden dimension**
 - Different **embedding**
 - Different **model**
 - Different type of model → **MLP, RNN, GNN...**
 - Maybe the different can identify the different single triplet class

Ensemble - Voting

- Noticed that **T1518_c9b** always got predicted in all the experiments
 - Euni's MLP and RNN also predict **T1518_c9b** perfectly → **Reason ?!**
- Experiment result on the single triplet case:

accuracy			0.02	5300
macro avg	0.02	0.02	0.02	5300
weighted avg	0.03	0.02	0.02	5300

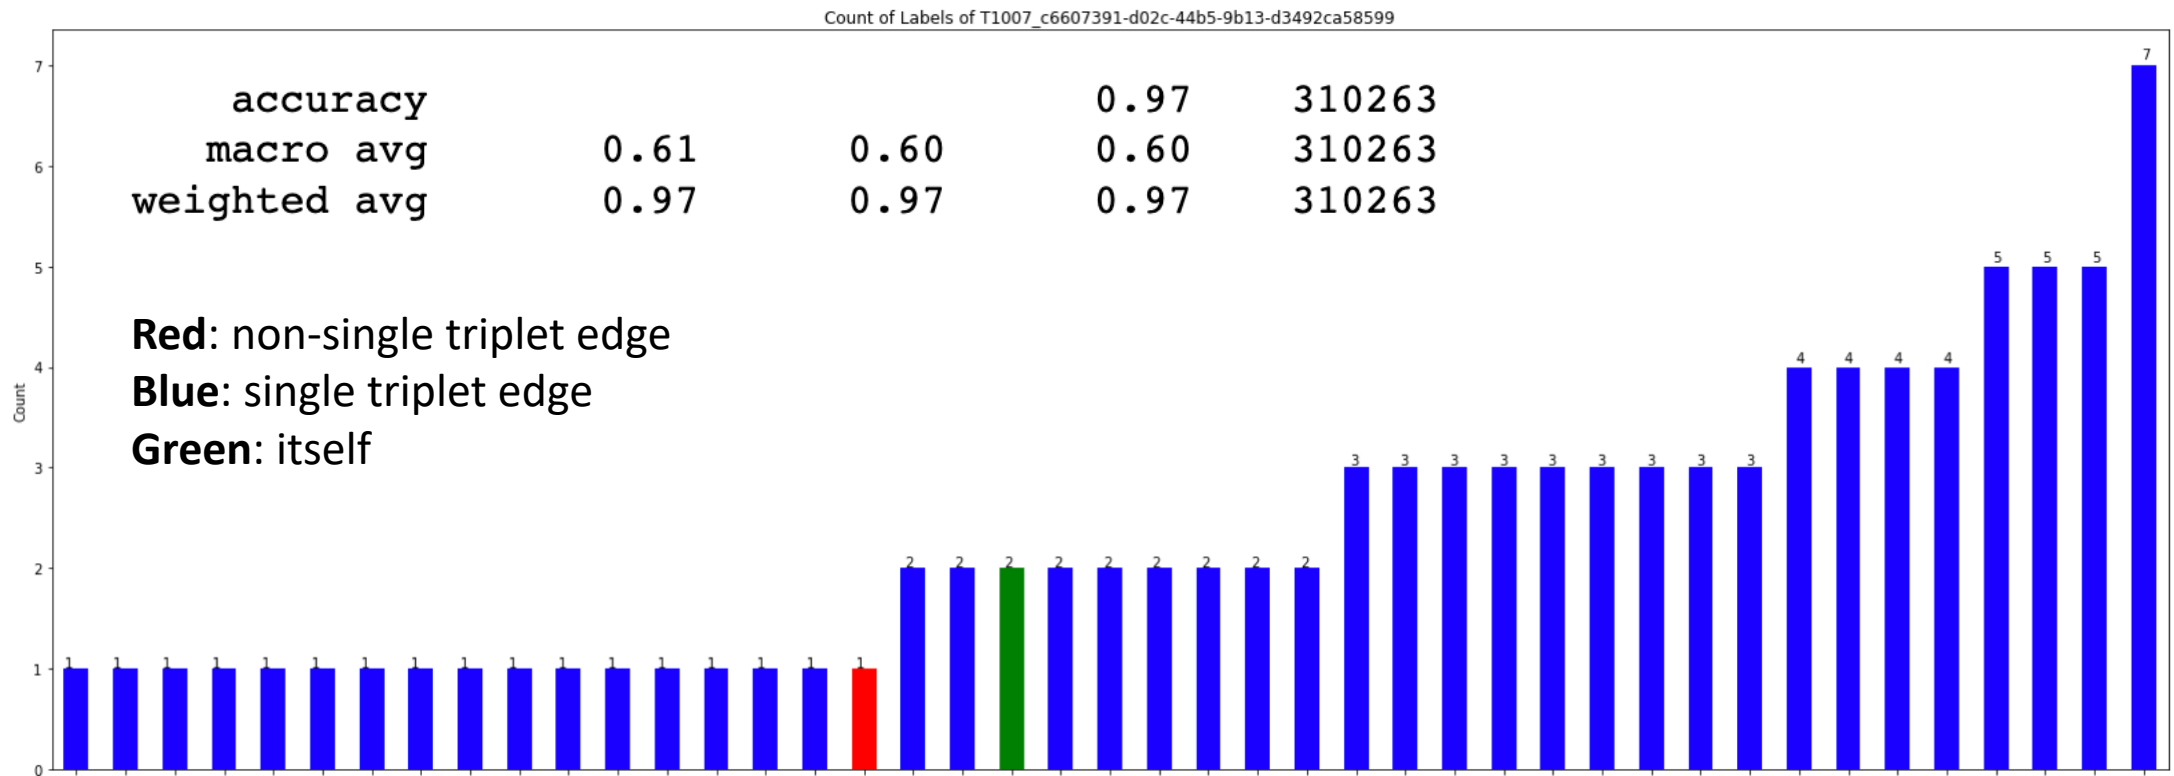
- Voting these three models (MLP, RNN, GNN) may not be useful

true_label	predicted_label
T1003.003_9f73269695e54311dd61dc68940fb3e1	T1490_9e5e4c0655fd1b5be88bd40b8251175f
T1003.003_9f73269695e54311dd61dc68940fb3e1	T1490_9e5e4c0655fd1b5be88bd40b8251175f
T1003.003_9f73269695e54311dd61dc68940fb3e1	T1490_9e5e4c0655fd1b5be88bd40b8251175f
T1003.003_9f73269695e54311dd61dc68940fb3e1	T1564.003_9a2edad4053a2b59fb9167a9bc29e7dc
T1003.003_9f73269695e54311dd61dc68940fb3e1	T1499_2fe2d5e6-7b06-4fc0-bf71-6966a1226731
T1003.003_9f73269695e54311dd61dc68940fb3e1	T1490_9e5e4c0655fd1b5be88bd40b8251175f

Since the order of the prediction would be the same

Observation of the Prediction

- Use the original training set
- The distribution of the prediction is so sparse
 - Most of the predictions are like this:



Observation of the Prediction

- The # of the predicted labels

Top 5 Labels and Counts:

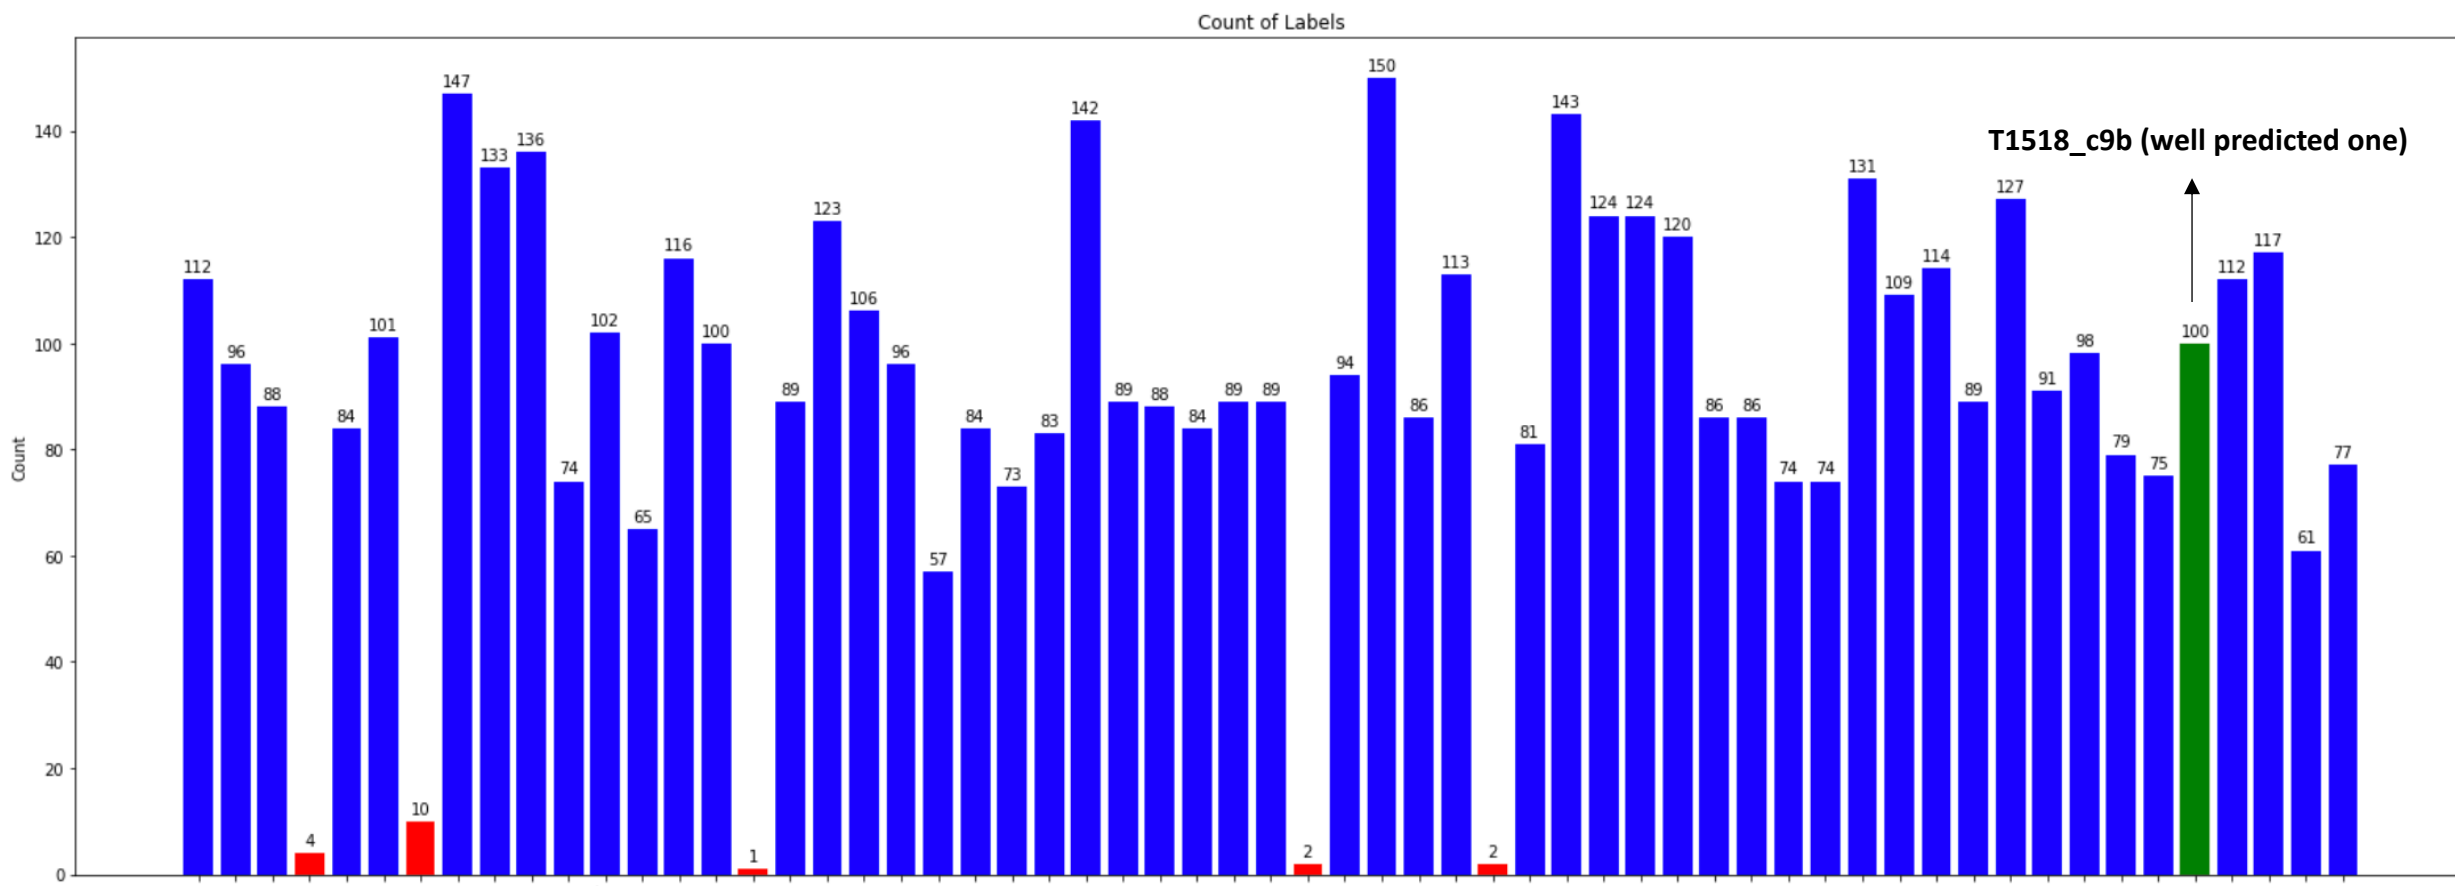
T1082_29451844-9b76-4e16-a9ee-d6feab4b24db: 150

T1016_921055f4-5970-4707-909e-62f594234d91: 147

T1124_fa6e8607-e0b1-425d-8924-9b894da5a002: 143

T1053.005_ee454be9197890de62705ce6255933fd: 142

T1016_e8017c46-acb8-400c-a4b5-b3362b5b5baa: 136



Experiment - Remove

Remove the Popular TTPs

Top 5 Labels and Counts:

T1082 29451844-9b76-4e16-a9ee-d6feab4b24db: 150

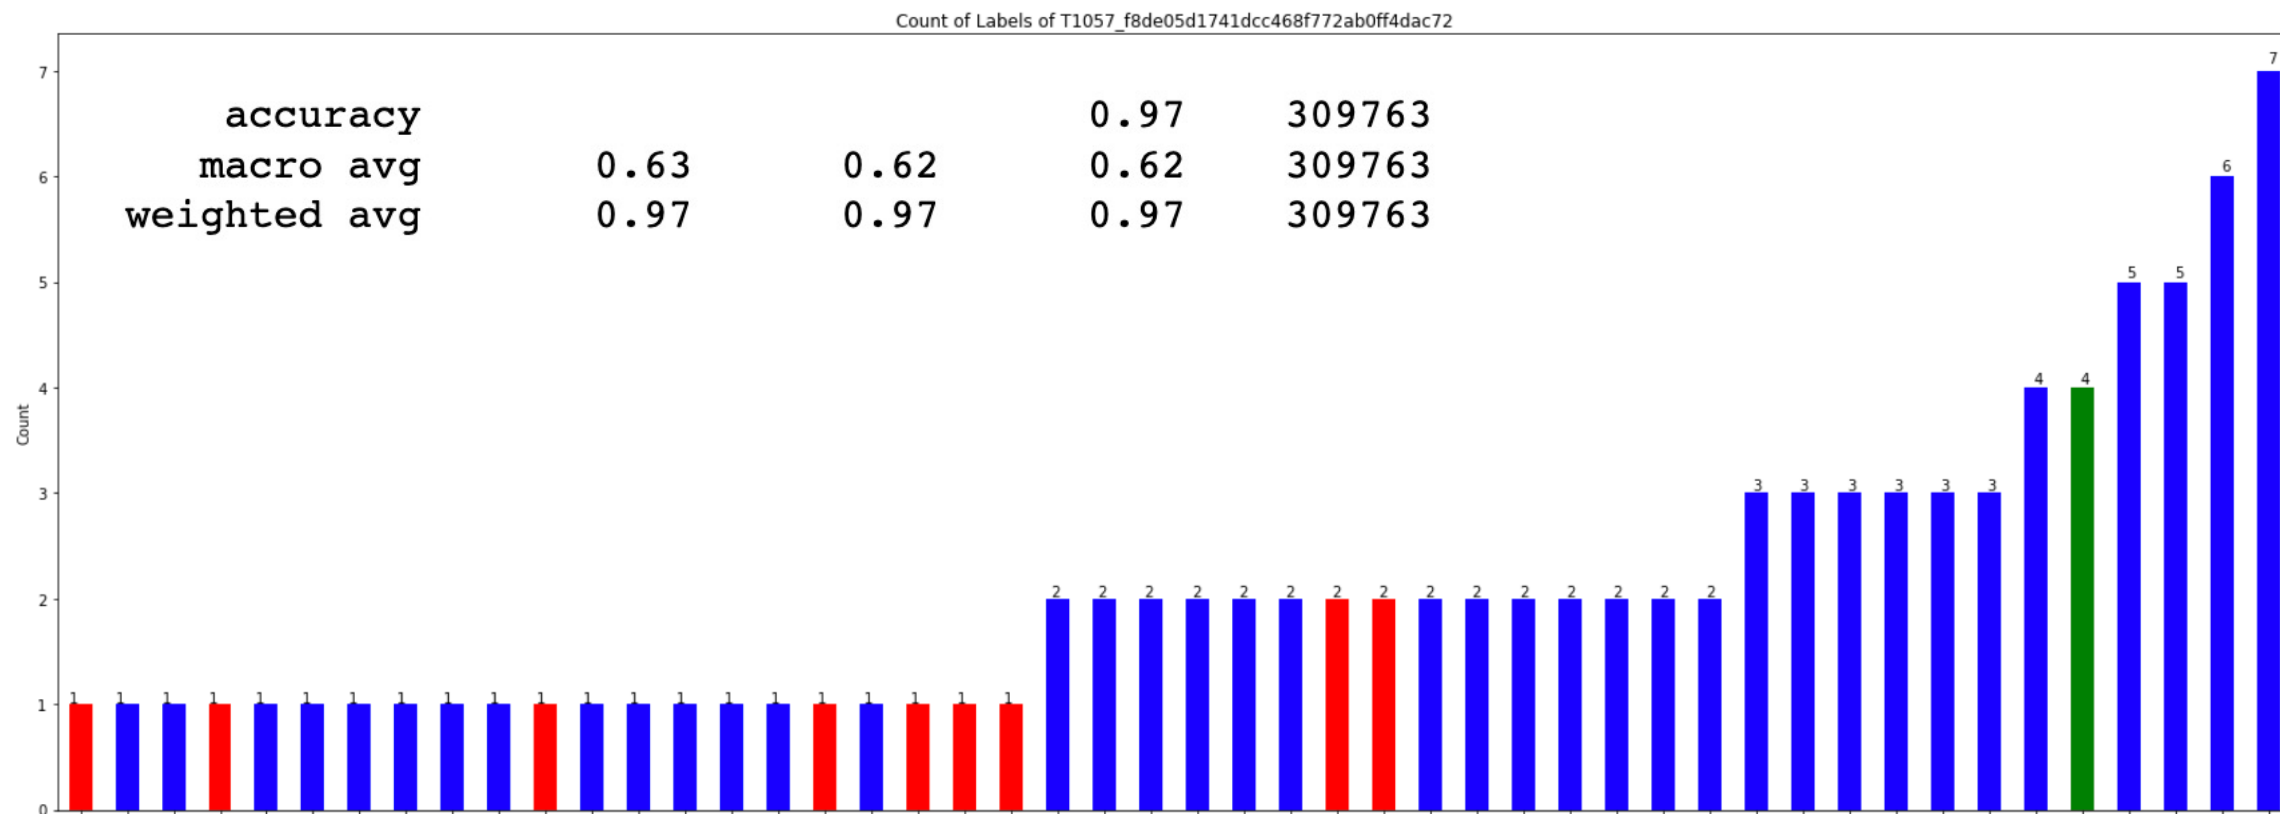
T1016 921055f4-5970-4707-909e-62f594234d91: 147

T1124 fa6e8607-e0b1-425d-8924-9b894da5a002: 143

T1053.005 ee454be9197890de62705ce6255933fd: 142

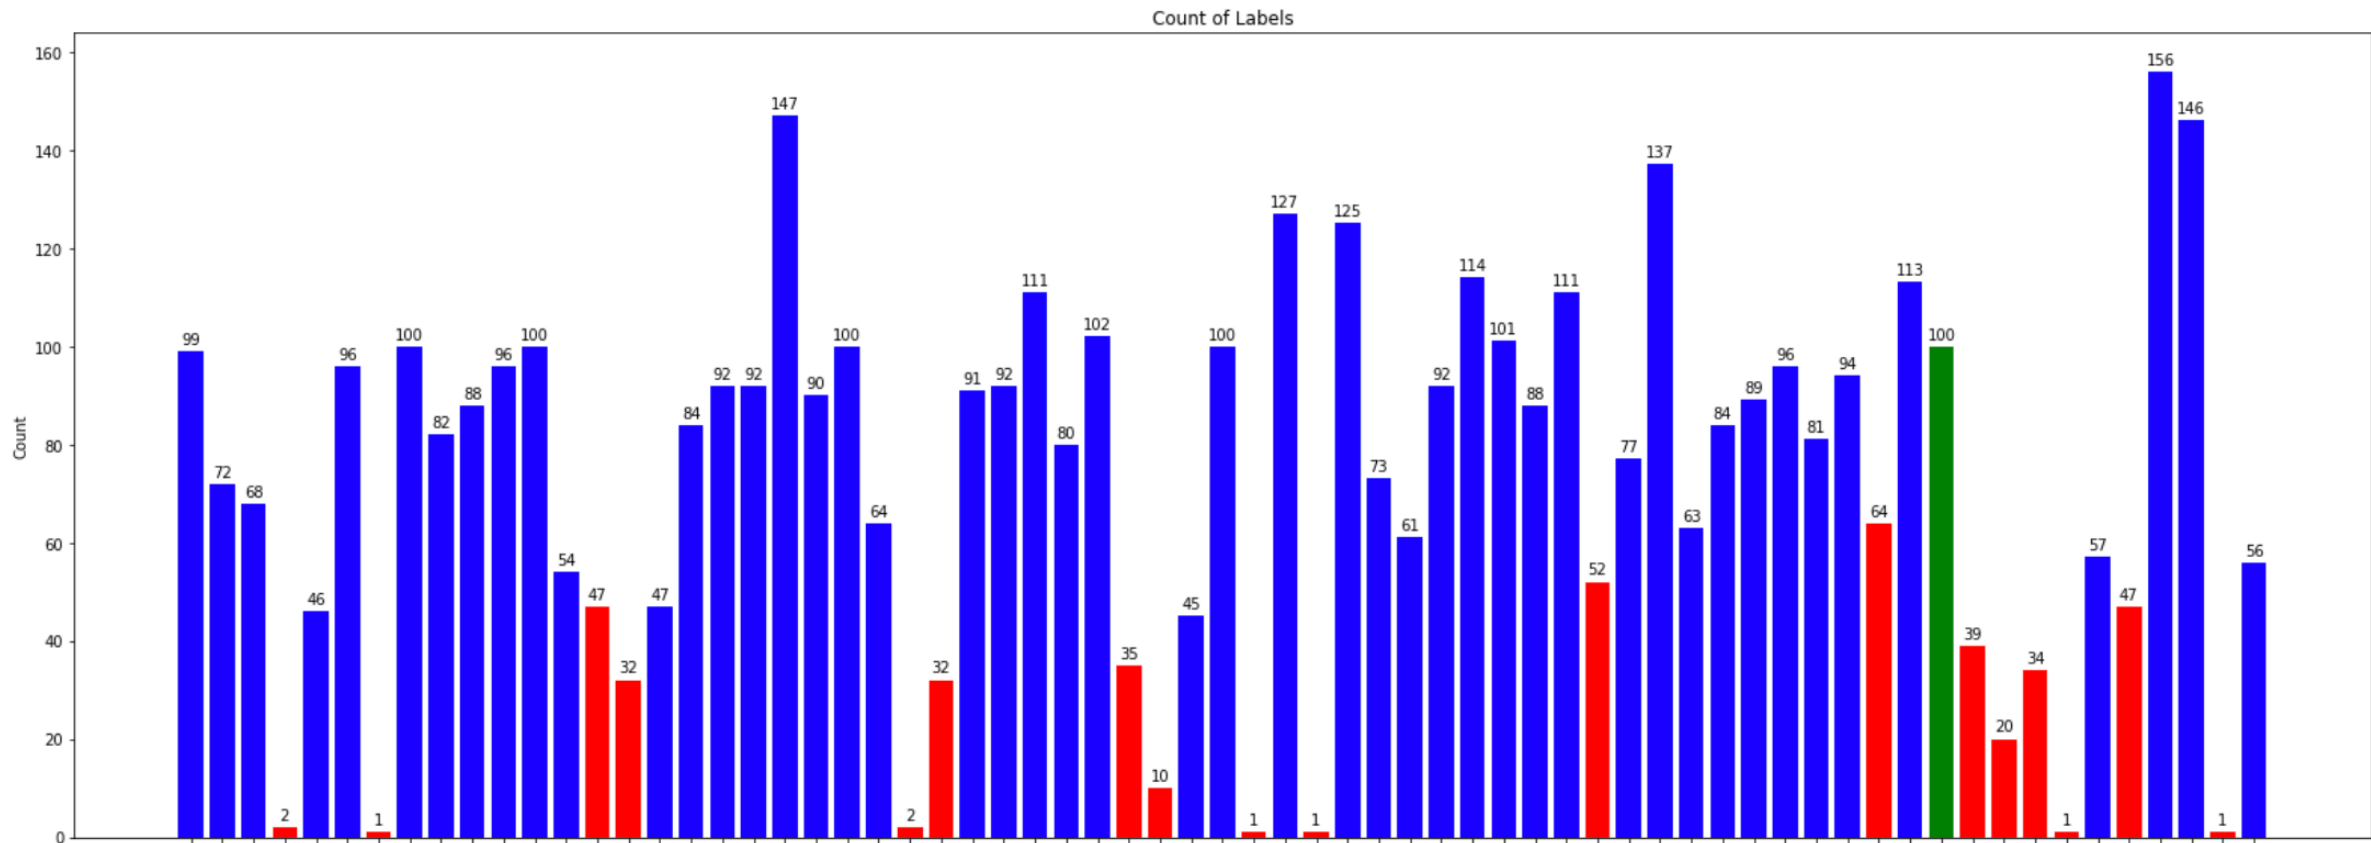
T1016 e8017c46-acb8-400c-a4b5-b3362b5b5baa: 136

- Remove these 5 TTPs from the dataset and then train it again:
 - More prediction on the non-single triplet case is like this:



Remove the Popular TTPs

- Remove these 5 TTPs from the dataset and then train it again:
- The # of the predicted labels



GNN - Paper

GraphSMOTE: Imbalanced Node Classification on Graphs with Graph Neural Networks

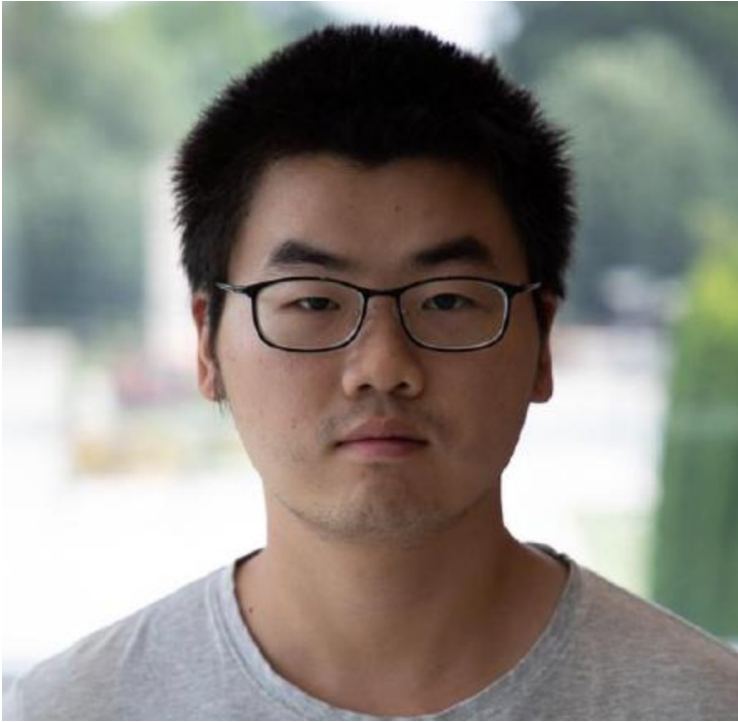
WSDM '21, March 8–12, 2021, Virtual Event, Israel

Tianxiang Zhao, Xiang Zhang, Suhang Wang

{tkz5084,xzz89,szw494}@psu.edu

College of Information Science and Technology, Penn State University State College, The USA

<https://github.com/TianxiangZhao/GraphSmote>

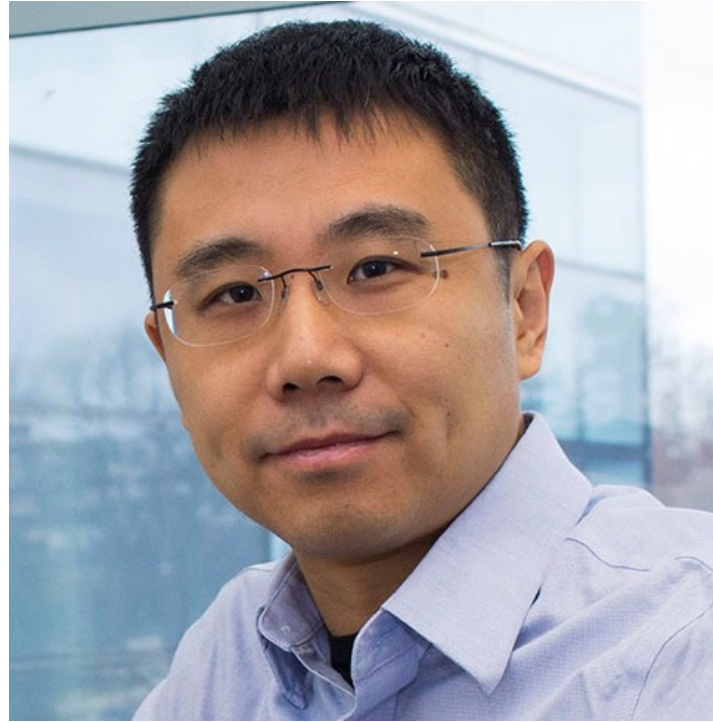


Tianxiang Zhao

PhD student in PSU

co-advised by Prof.Xiang Zhang and
Prof.Suhang Wang.

<https://scholar.google.com/citations?user=pXkPq3YAAAAJ&hl=en>



Xiang Zhang

Associate Professor in PSU

Data Mining, Machine Learning and
Big Data Analytics

<https://ist.psu.edu/directory/xzz89>



Suhang Wang

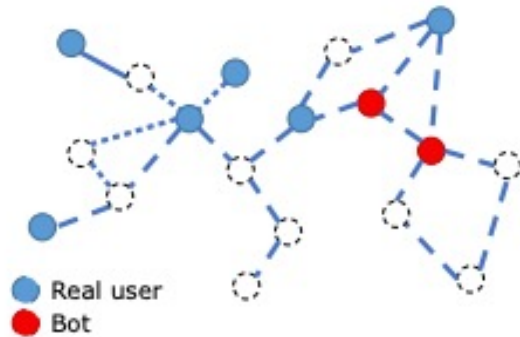
Associate Professor in PSU

Data Mining, Machine Learning and
Social Media Mining

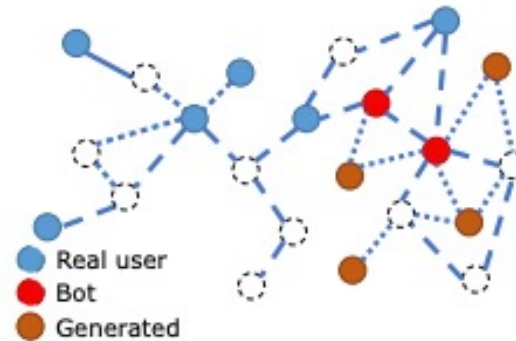
<https://suhangwang.ist.psu.edu/>

GraphSMOTE

- Task:

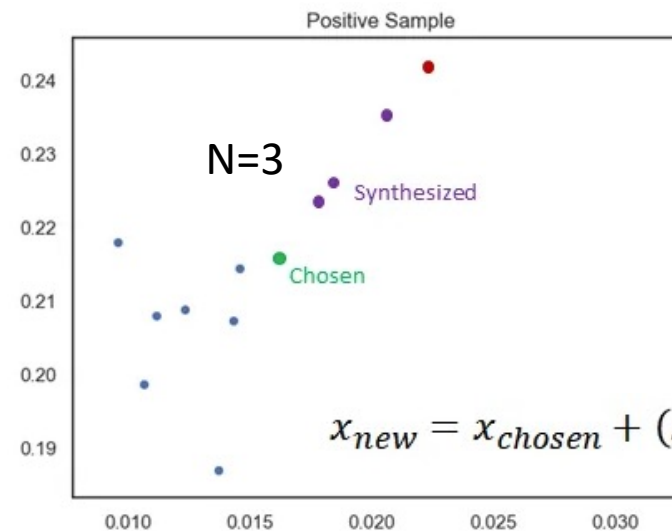
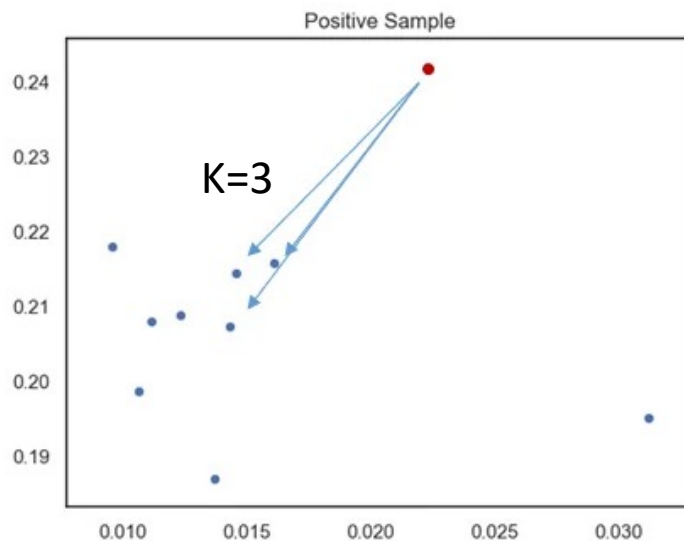


(a) Bot detection task



(b) After over-sampling

- Synthesized Minority Oversampling Technique (SMOTE)



- Used on i.i.d. data
- Doesn't consider the structure of the graph
- No edge connection between the nodes

GraphSMOTE

- **Framework**

- a GNN-based feature extractor
→ Use **GraphSAGE**
- Synthetic Node Generation
→ Use **SMOTE** algorithm
- Edge Generator
→ weighted inner product **decoder** F
- GNN Classifier (downstream task)

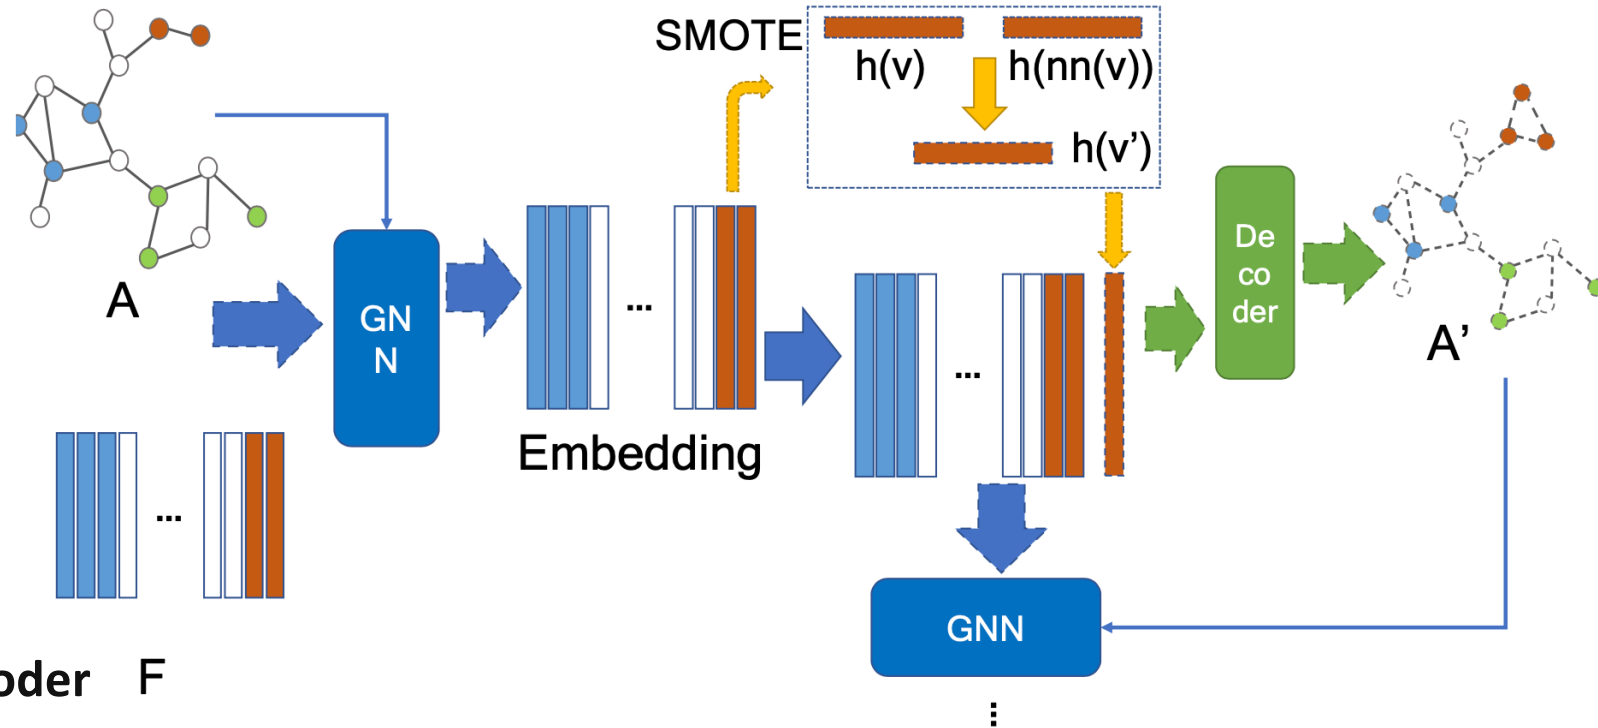


Figure 2: Overview of the framework

Future Work

Future Work

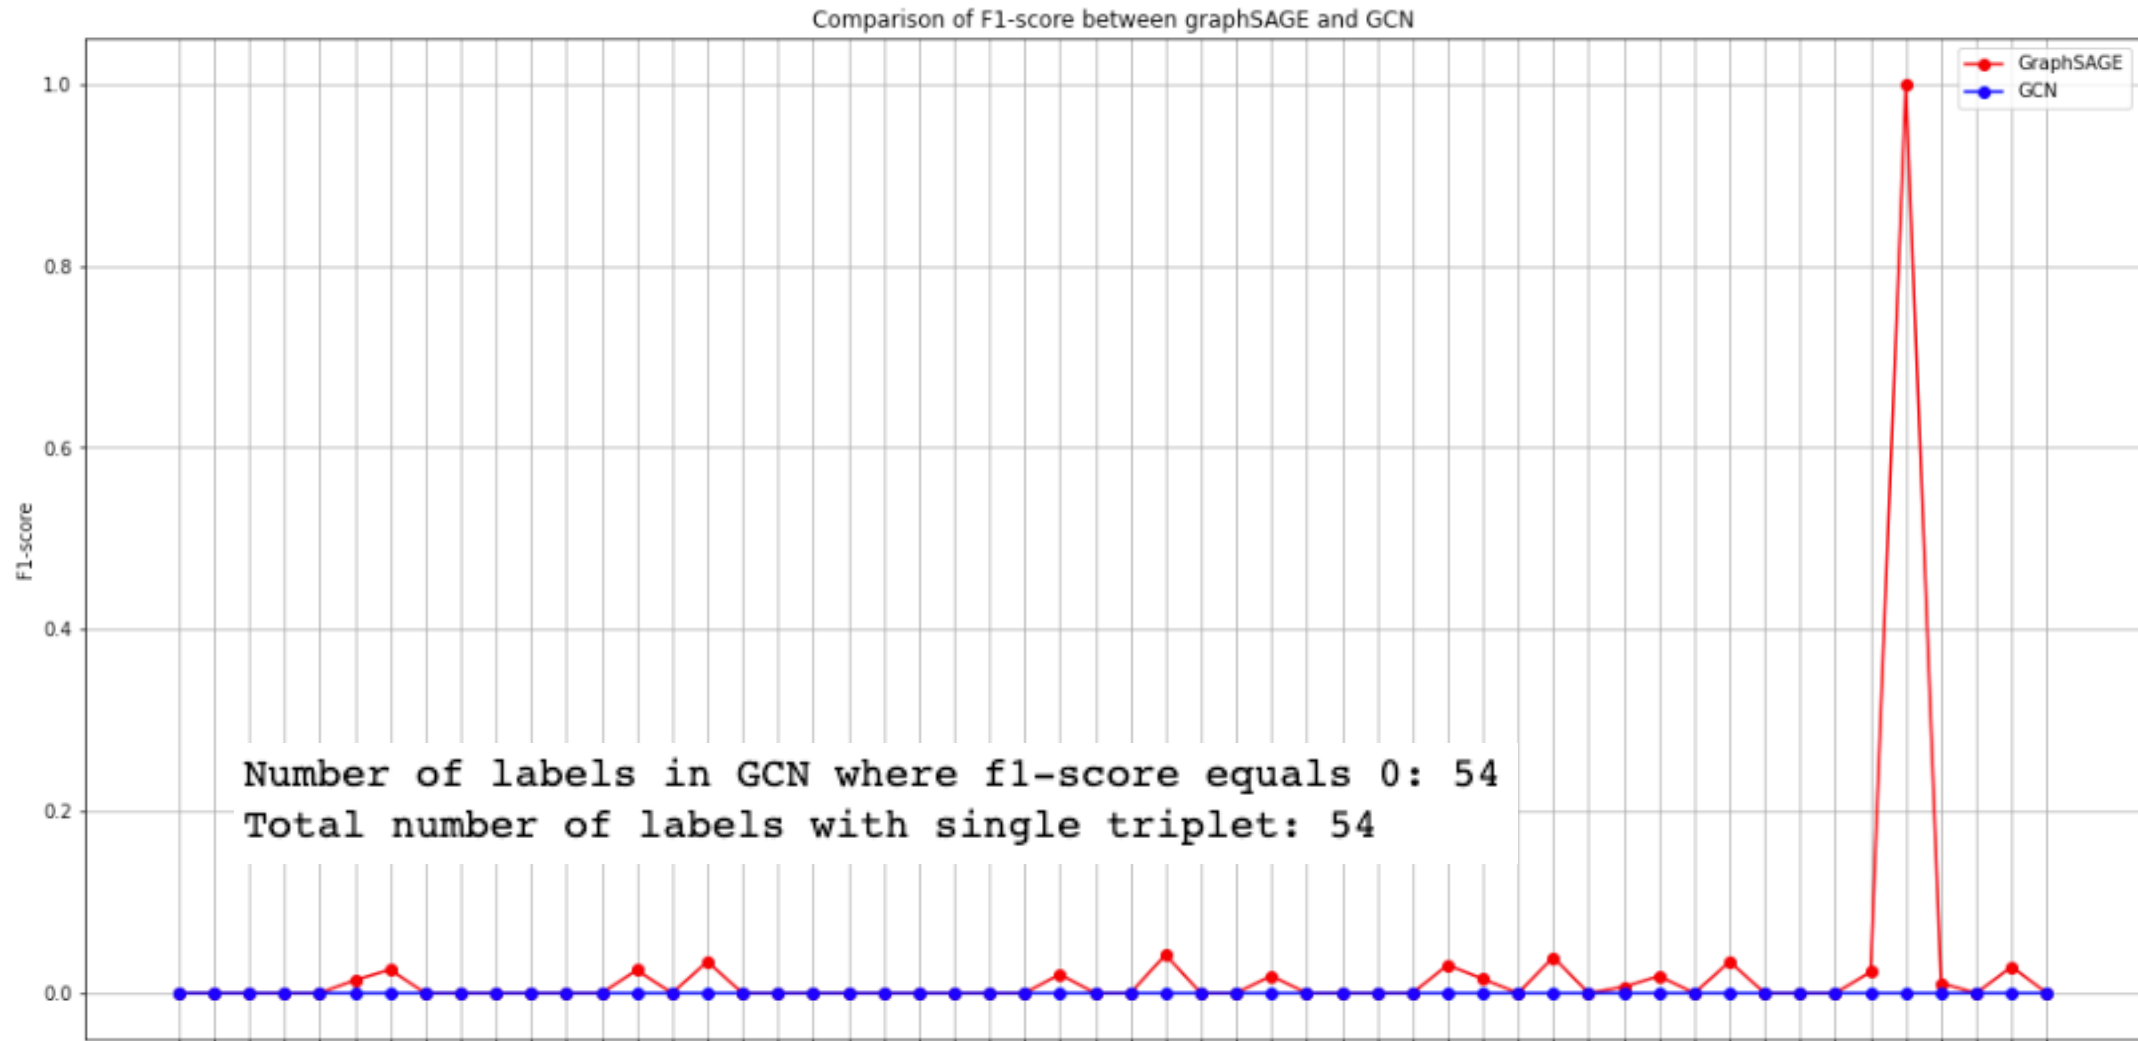
- **GNN**

- Try some other methods to improve the performance of single triplet issue
 - Figure out why the model can detect the **T1518_c9b** (if available?)
 - Read the **GraphSMOTE** paper

Thanks!!

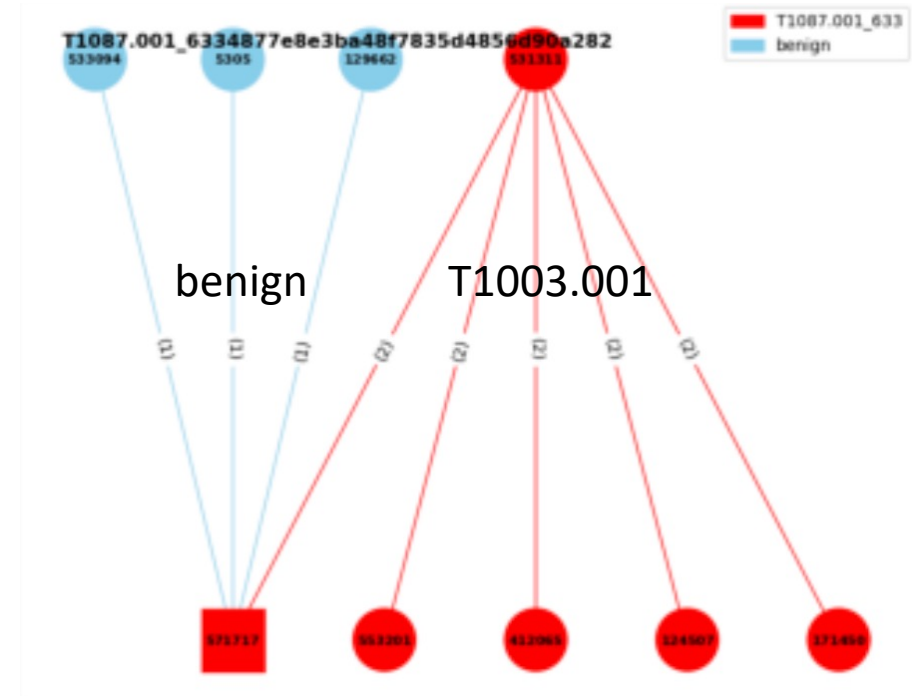
Appendix

Observation on Different Model



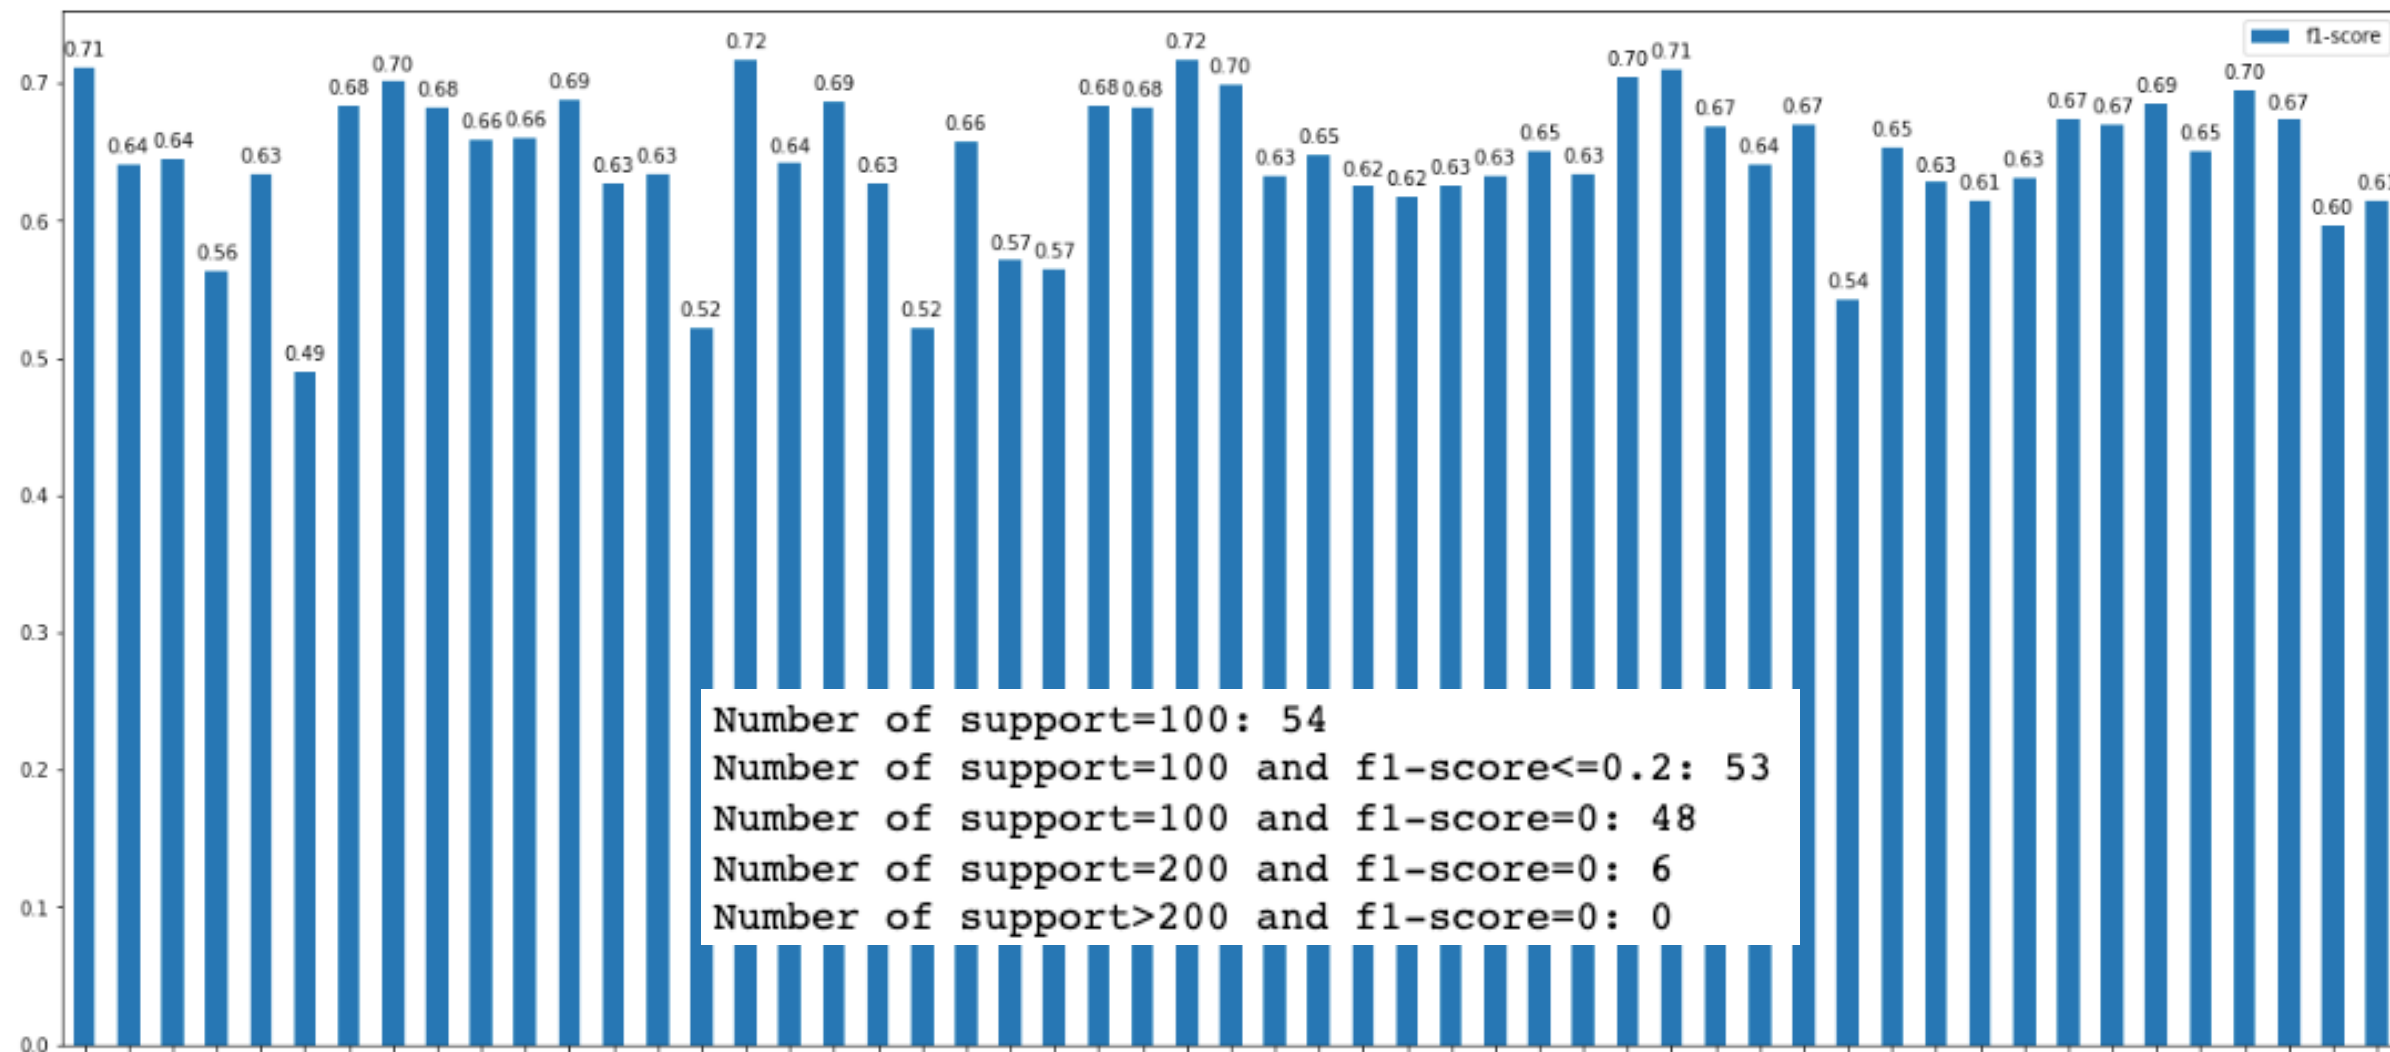
Experiment 3

- **Experiment 3:**
 - Consider the **neighbor** benign nodes
 - Edge classification
- Given a graph \rightarrow label the triplets with the benign or the specific AP



Experiment 3 - Oversampling

Bar Chart of hidden dimension = 256:

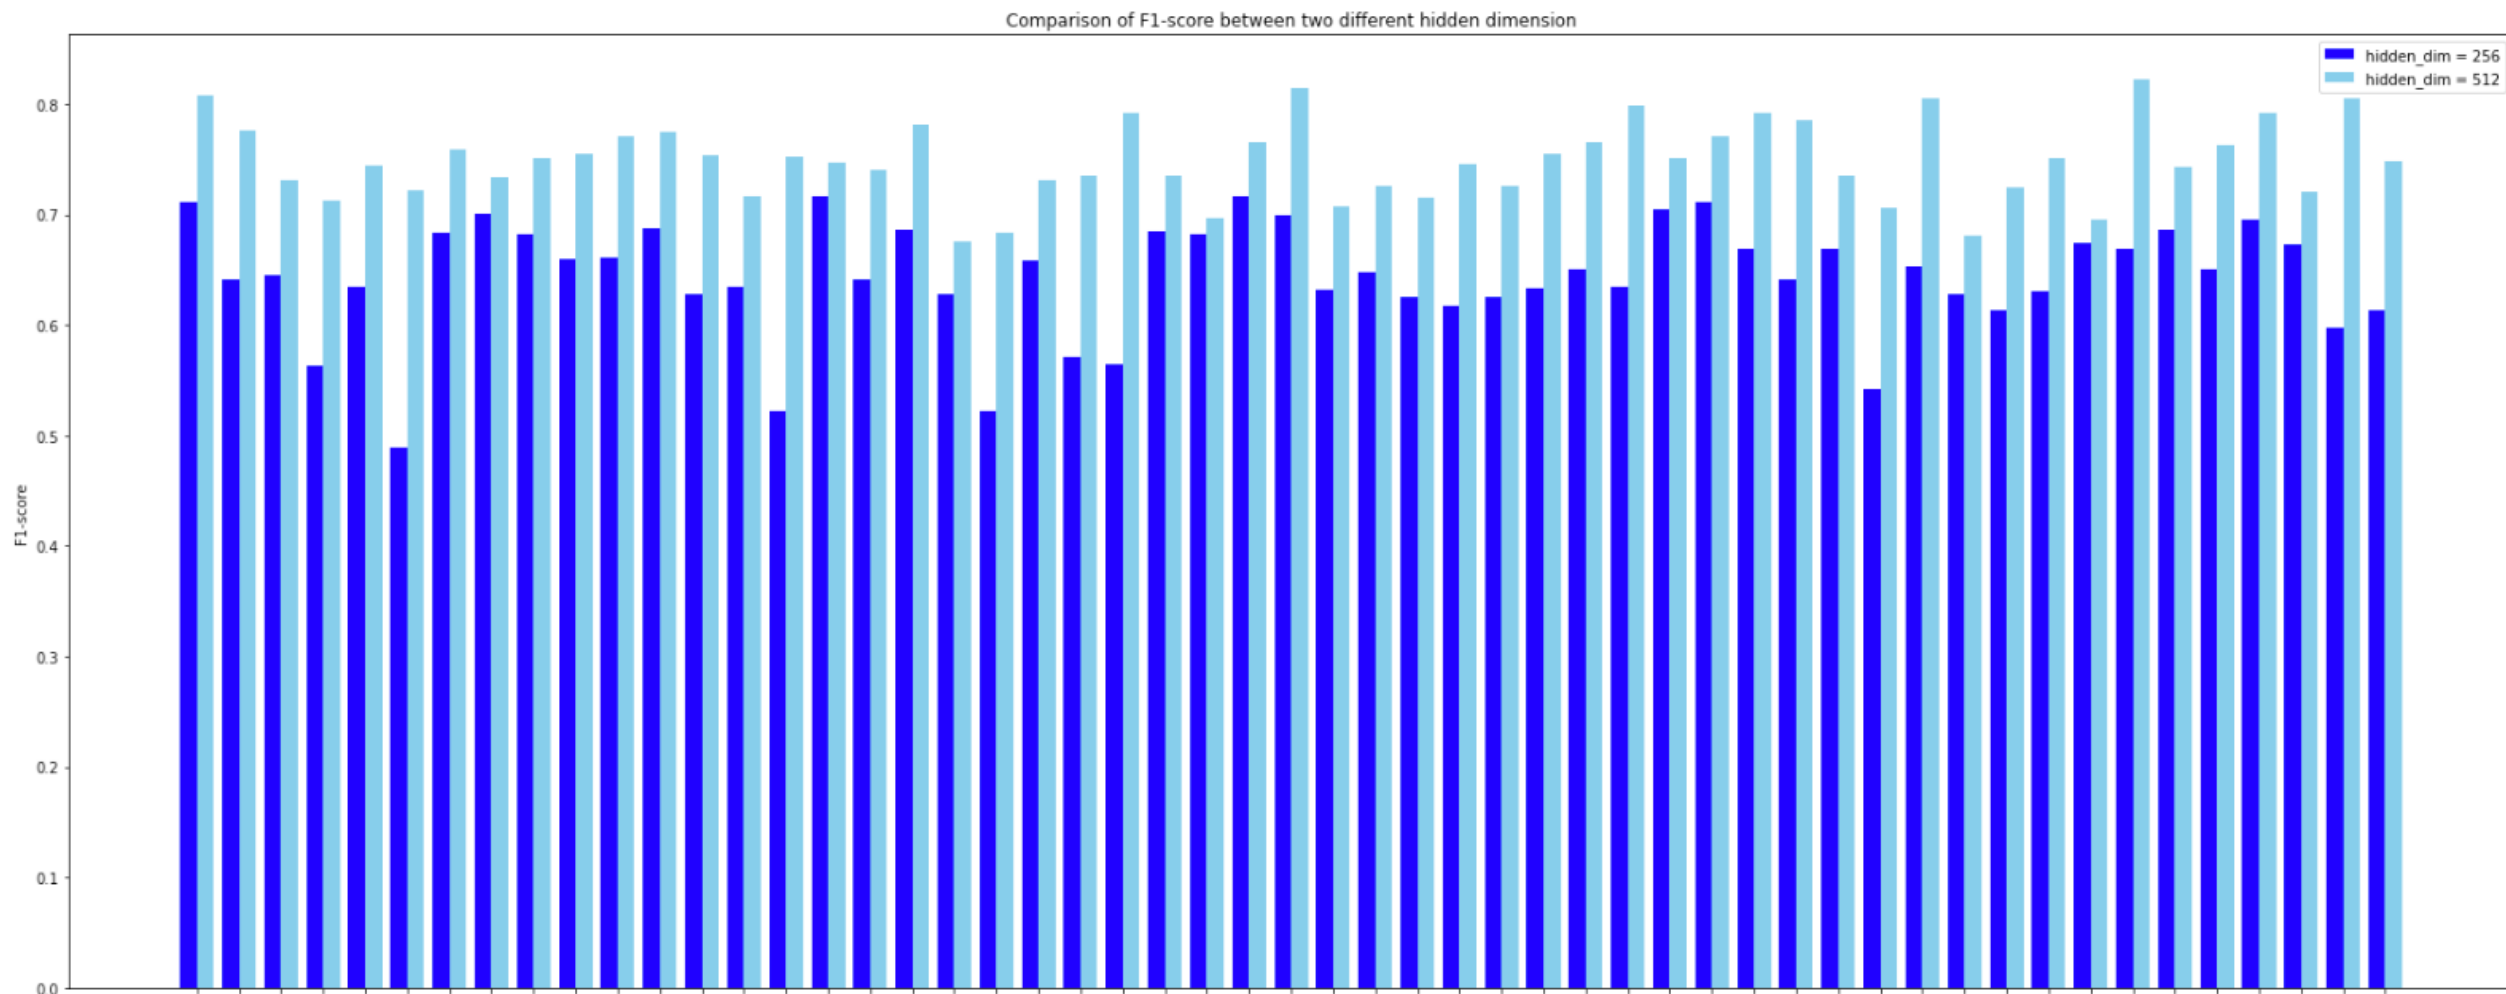


Experiment 3 - Oversampling

- Current Trial: Duplicate the data with single triplets → 20, 40, 80, 320 times

20 times	Number of support=100: 54				
	Number of support=100 and f1-score<=0.2: 53				
	Number of support=100 and f1-score=0: 21				
	Number of support=200 and f1-score=0: 10				
	Number of support>200 and f1-score=0: 1				
40 times	Number of support=100: 54	accuracy	0.971560	0.971560	0.971560
	Number of support=100 and f1-score<=0.2: 53				0.97156
	Number of support=100 and f1-score=0: 14	macro avg	0.597445	0.600577	0.594684
	Number of support=200 and f1-score=0: 10				310263.00000
	Number of support>200 and f1-score=0: 2	weighted avg	0.970613	0.971560	0.970482
80 times	Number of support=100: 54				310263.00000
	Number of support=100 and f1-score<=0.2: 53	accuracy	0.971318	0.971318	0.971318
	Number of support=100 and f1-score=0: 18				0.971318
	Number of support=200 and f1-score=0: 10	macro avg	0.602542	0.597866	0.594387
	Number of support>200 and f1-score=0: 3				310263.00000
		weighted avg	0.971349	0.971318	0.970477
					310263.00000
		accuracy	0.971463	0.971463	0.971463
					0.971463
		macro avg	0.596490	0.598077	0.594237
					310263.00000
		weighted avg	0.970626	0.971463	0.970567
					310263.00000

Experiment 3 - Oversampling



Experiment 3 - Model

- Concept from the DGL official website:
 1. Let the dgl graph's edge data have the attribute: **edata["label"]**
 2. Use **GraphSAGE** model to get the new **node embedding**
 3. Use **MLP** model to get the **score** of the edge
 4. Concatenate these two models
 5. Train the final model

```
g.ndata['feat'] = th.tensor(data["node_feat"])
g.edata['feat'] = th.tensor(data["edge_attr"])
g.edata['label'] = th.tensor(data["labels"])
```

```
def model_fn(batched_g, model, criterion, device, count=1, which_type='train'):
    """Forward a batch through the model."""
    batched_g = batched_g.to(device)
    labels = batched_g.edata['label'].to(device)

    logits = model(batched_g, batched_g.ndata['feat'].float())
    loss = criterion(logits, labels)

    output = torch.softmax(logits, dim=1)
    preds = output.argmax(1)

    accuracy = torch.mean((preds == labels).float())
```


Experiment 3 - Model

```
class GraphSAGE(nn.Module):
    def __init__(self, in_dim, hidden_dim, out_dim):
        super(GraphSAGE, self).__init__()
        self.layer1 = dgl.nn.SAGEConv(in_dim, hidden_dim, 'pool')
        self.layer2 = dgl.nn.SAGEConv(hidden_dim, out_dim, 'pool')

    def forward(self, g, inputs):
        h = self.layer1(g, inputs)
        h = torch.relu(h)
        h = self.layer2(g, h)
        return h
```

```
class MLPPredictor(nn.Module):
    def __init__(self, out_feats, out_classes):
        super().__init__()
        self.W = nn.Linear(out_feats*2, out_classes)

    def apply_edges(self, edges):
        h_u = edges.src['h']
        h_v = edges.dst['h']
        score = self.W(torch.cat([h_u, h_v], 1))
        return {'score': score}

    def forward(self, graph, h):
        with graph.local_scope():
            graph.ndata['h'] = h
            graph.apply_edges(self.apply_edges)
            return graph.edata['score']
```

```
class Model(nn.Module):
    def __init__(self, in_features, hidden_features, out_features, num_classes):
        super().__init__()
        self.sage = GraphSAGE(in_features, hidden_features, out_features)
        self.pred = MLPPredictor(out_features, num_classes)

    def forward(self, g, node_feat, return_logits=False):
        h = self.sage(g, node_feat)
        logits = self.pred(g, h)

        return logits
```

Experiment 3 - Result

- **Format of the edge labels:**

- Label 65 is benign

```
labels of Test: tensor([155, 65, 155, 155, 155], device='cuda:0') torch.Size([5])
predicted of Test: tensor([155, 65, 155, 155, 155], device='cuda:0') torch.Size([5])
labels of Test: tensor([61, 61, 61], device='cuda:0') torch.Size([3])
predicted of Test: tensor([61, 61, 61], device='cuda:0') torch.Size([3])
```

- **Classification report:**

- transR_50:

4a0dc2e1f5d1a	0.00	0.00	0.00	100
167175e8a019a	1.00	1.00	1.00	800
c3579e9e3737b	1.00	1.00	1.00	6200
43d838e0791ca	1.00	1.00	1.00	600
benign	1.00	1.00	1.00	134563
accuracy			0.97	310263
macro avg	0.60	0.61	0.60	310263
weighted avg	0.97	0.97	0.97	310263

- secureBERT_50:

714a0dc2e1f5d1a	0.00	0.00	0.00	100
fb167175e8a019a	0.98	1.00	0.99	800
2ac3579e9e3737b	0.97	0.98	0.98	6200
0243d838e0791ca	0.91	0.83	0.87	600
benign	0.99	1.00	0.99	134563
accuracy			0.92	310263
macro avg	0.52	0.48	0.49	310263
weighted avg	0.90	0.92	0.91	310263

- **Macro average** is similar to previous experiments → won't be affected by benign
- **Weighted average** is very high since the # of the benign is high(unbalanced) and predictable
- TransX family performs better than secureBERT

Experiment 3 – Noise

- Current Trial 1:
 - Add the **noise** to the node feature

```
def collate(samples):
    data_list = samples
    batched_graphs = []
    for data in data_list:
        g = dgl.graph((th.tensor(data["edge_index"])[0]), th.tensor(data["edge_index"])[1]), num_nodes=data["num_nodes"])

        node_feat = th.tensor(data["node_feat"])

        noise = th.normal(mean=0, std=0.01, size=node_feat.shape, device=node_feat.device)
        node_feat += noise

        g.ndata['feat'] = node_feat
        g.edata['feat'] = th.tensor(data["edge_attr"])
        g.edata['label'] = th.tensor(data["labels"]) # Add edge labels to graph

        batched_graphs.append(g)

    return dgl.batch(batched_graphs)
```

```
Number of support=100: 54
Number of support=100 and f1-score<=0.2: 53
Number of support=100 and f1-score=0: 46
Number of support=200 and f1-score=0: 4
Number of support>200 and f1-score=0: 0
```

Experiment 3 – K-fold validation

k-fold cross-validation (k 折交叉驗證) 是一種在機器學習中常用的模型評估方法，尤其在有限的數據集上評估模型性能時非常有效。其主要優點包括：

1. **更可靠的性能估計**：通過將數據集分成 k 個子集，每次使用其中一個子集作為測試集，其餘 $k-1$ 個子集作為訓練集，然後重複此過程 k 次，每次選擇不同的子集作為測試集，可以使得模型評估的結果更加穩定和可靠。
2. **充分利用數據**：每個數據點都被用作 $k-1$ 次的訓練和 1 次的測試，這意味著每個數據點都被充分利用，這在數據量不大時尤其重要。
3. **減少偏差**：由於模型需要在 k 個不同的訓練集上訓練，然後在 k 個不同的測試集上測試，這有助於降低模型在特定數據集上的性能估計偏差。
4. **更好的泛化性能評估**：通過對不同的訓練集和測試集進行訓練和測試，可以幫助評估模型對未見數據的泛化能力。
5. **減少過擬合風險**：k-fold cross-validation 有助於識別模型是否對特定的訓練數據集過度擬合，因為它必須在多個不同的訓練集上表現良好。

然而，k-fold cross-validation 也有其局限性，例如計算成本較高（尤其是 k 較大和模型訓練時間較長時），並且結果可能依賴於 k 的選擇以及數據分割的方式。這些都是在使用此方法時需要考慮的因素。